

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>

PROJECTS : BUILDING Residual Networks very deep networks. The problem with deep networks is vanishing gradients , but in RESnets we skip



deeplearning.ai

# Case Studies

---

Why look at  
case studies?

# Outline

## Classic networks:

- LeNet-5 ←
- AlexNet ←
- VGG ←

ResNet (152)

Inception



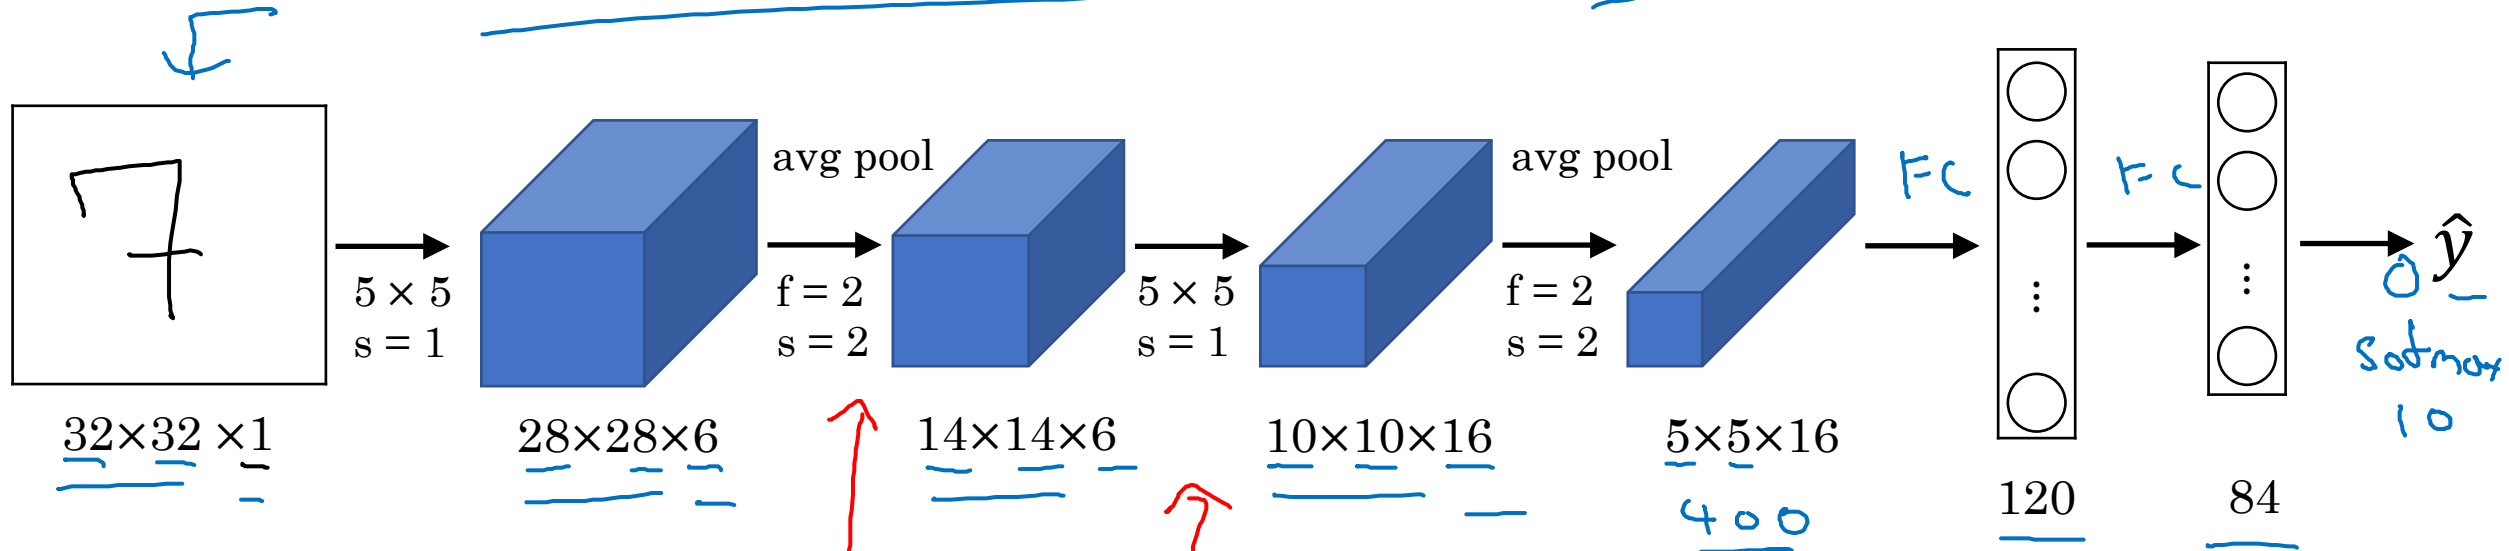
deeplearning.ai

# Case Studies

---

## Classic networks

# LeNet - 5



60K parameters.

$n_H, n_W \downarrow$   $n_C \uparrow$


conv pool conv pool fc fc output

Advanced: sigmoid/tanh ReLU

II, III.

↓

# AlexNet




- Similar to LeNet, but much bigger.

- ReLU

- Multiple GPUs.

- Local Response Normalization (LRN)




~60M parameters

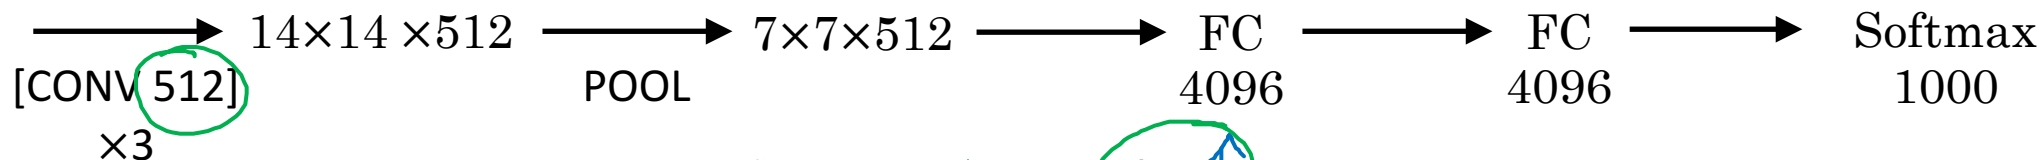
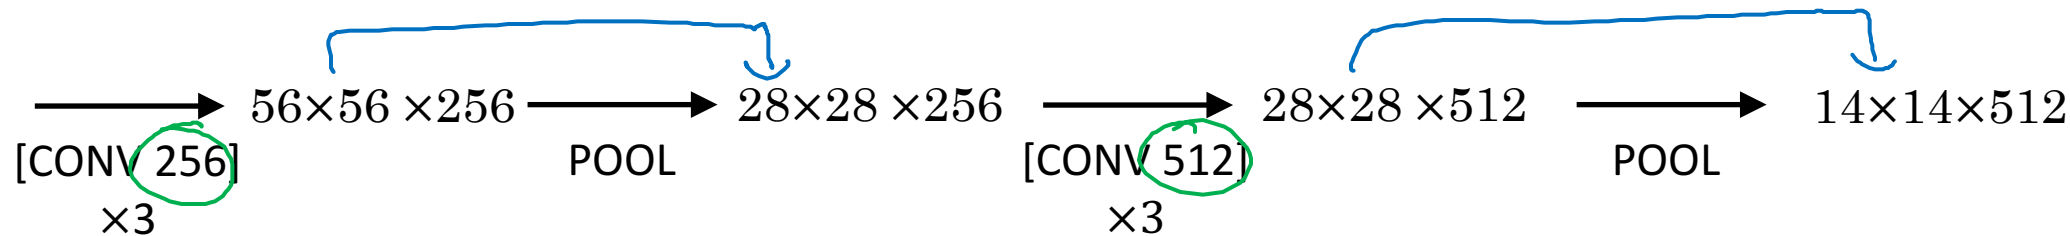
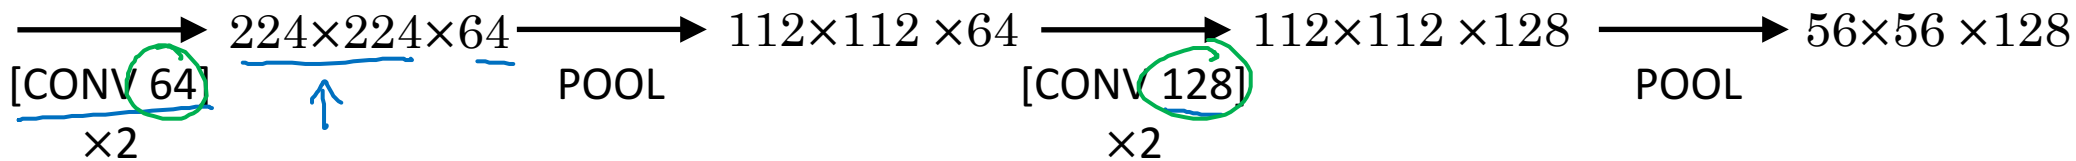
# VGG - 16

CONV = 3x3 filter, s = 1, same

MAX-POOL = 2x2, s = 2



224x224x3



$n_h, n_w \downarrow$

$n_c \uparrow$

$\sim 138M$

RESNETS allow us to train on much deeper network of over 100 layers. In classic network each information goes through all the layers , but in residua



deeplearning.ai

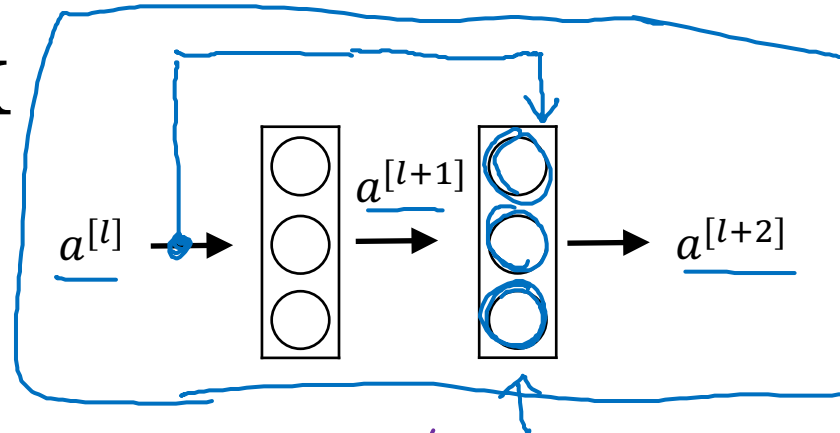
# Case Studies

---

## Residual Networks (ResNets)



# Residual block



"short cut" / skip connection



$$\underline{z^{[l+1]}} = W^{[l+1]} \underline{a^{[l]}} + b^{[l+1]}$$

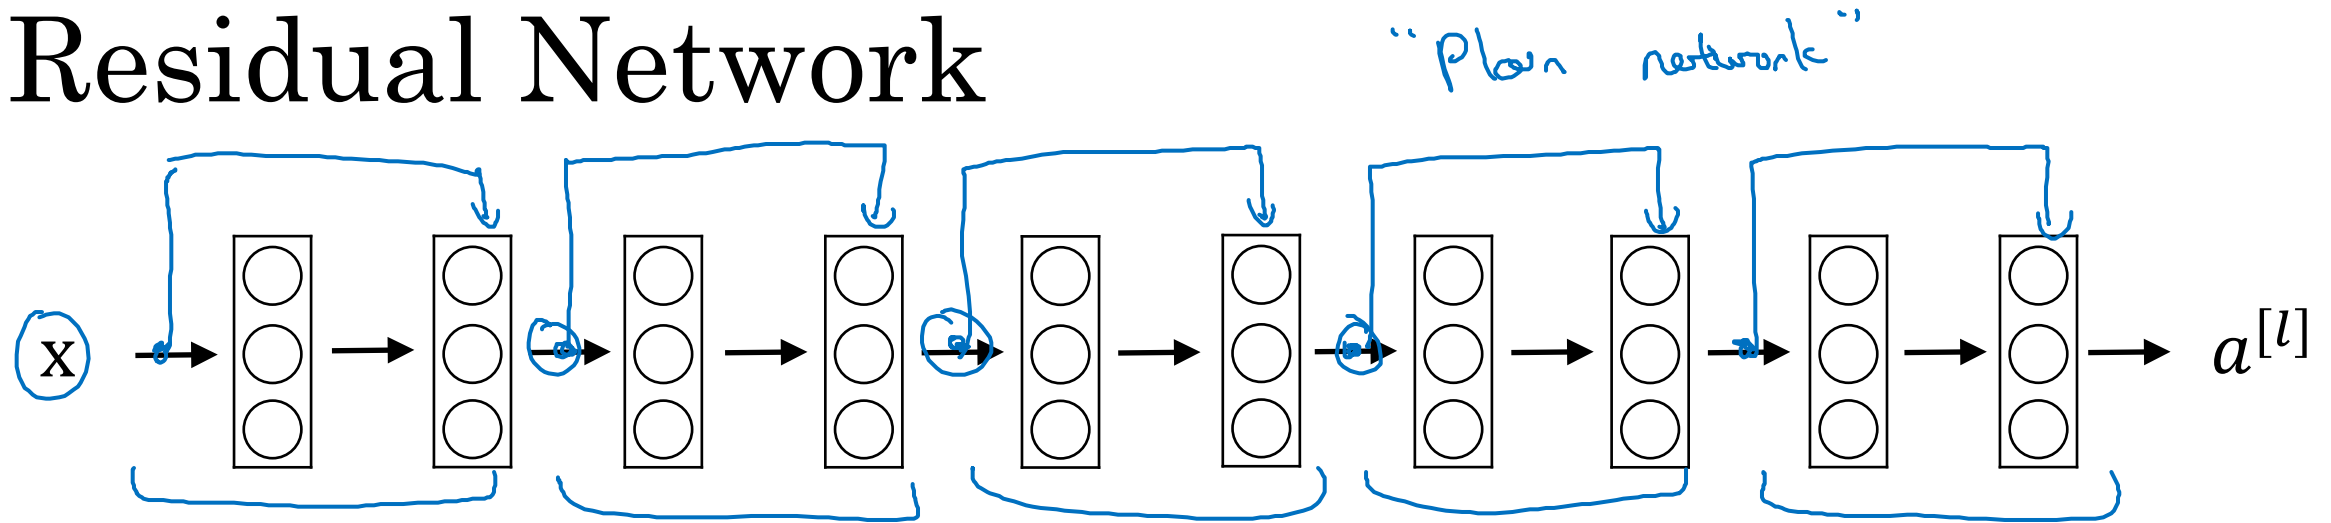
$$\underline{a^{[l+1]}} = g(\underline{z^{[l+1]}})$$

$$\underline{z^{[l+2]}} = W^{[l+2]} \underline{a^{[l+1]}} + b^{[l+2]}$$


~~$$a^{[l+2]} = g(z^{[l+2]})$$~~

$$a^{[l+2]} = g(z^{[l+2]} + \underline{a^{[l]}})$$


# Residual Network



Plain



ResNet



In theory increasing the depth of NN should continuously do more good on training set, but it's not the case in general. But this is true for RESNETs. RE



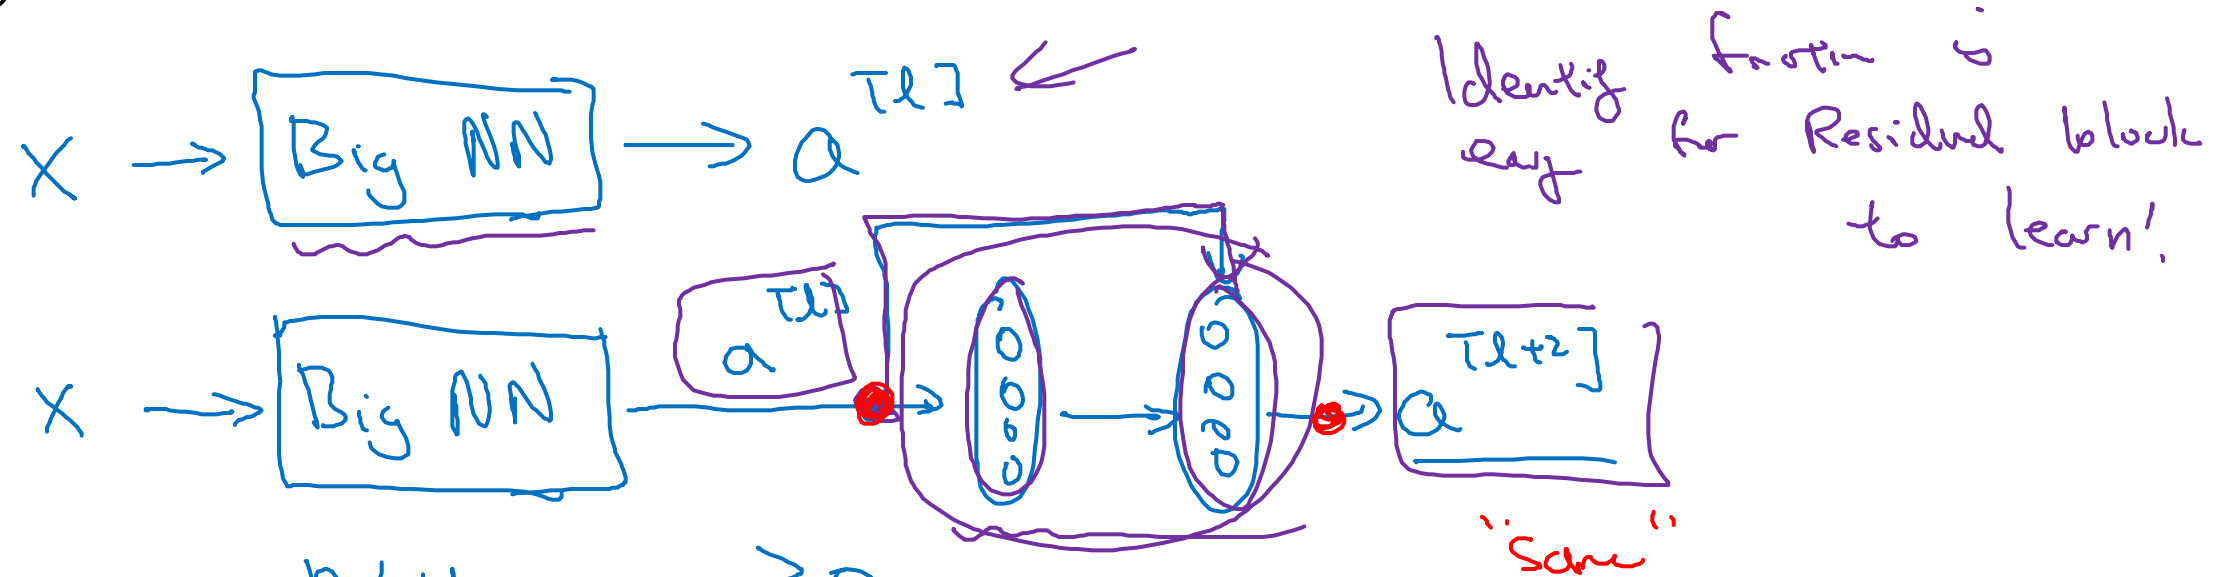
deeplearning.ai

# Case Studies

---

## Why ResNets work

# Why do residual networks work?



ReLU.  $a \geq 0$

$$a^{[L+2]} = g(\underbrace{z^{[L+2]} + a^{[L]}}_{\text{ReLU input}})$$

$$= g(\underbrace{W^{[L+2]} a^{[L]} + b^{[L+2]}}_{\text{if } W^{[L+2]}=0, b^{[L+2]}=0} + \underbrace{W_s a^{[L]}}_{\text{ReLU input}})$$

$256 \times 128$

$256 \times 128$

$128$

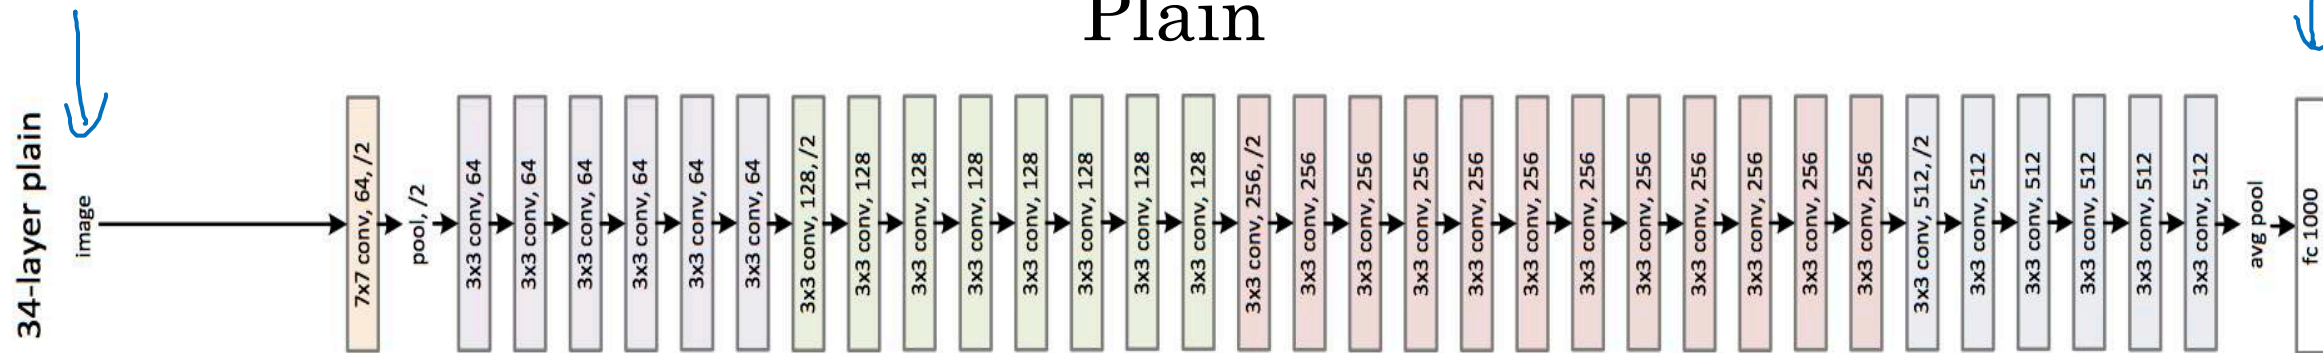
$$= g(a^{[L]})$$

$$= \underline{a^{[L]}}$$


# ResNet

dimension of  $a[l+2]$  should be same as that of  $a[l]$  so same convolution is generally performed. and If not happen we add a

## Plain



## ResNet





deeplearning.ai

# Case Studies


---

## Network in Network and $1\times 1$ convolutions

# Why does a $1 \times 1$ convolution do?

1	2	3	6	5	8
3	5	5	1	3	4
2	1	3	4	9	3
4	7	8	5	7	9
1	5	3	7	4	8
5	4	9	8	3	5

$6 \times 6 \times 1$




$6 \times 6 \times 32$

\*

2

=

32  $\rightarrow$  # filters.  
 $n_c^{[l+1]}$



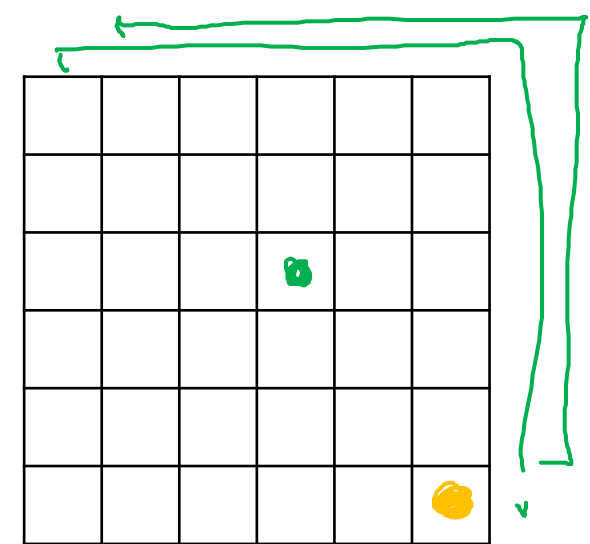
$1 \times 1 \times 32$

=

ReLU

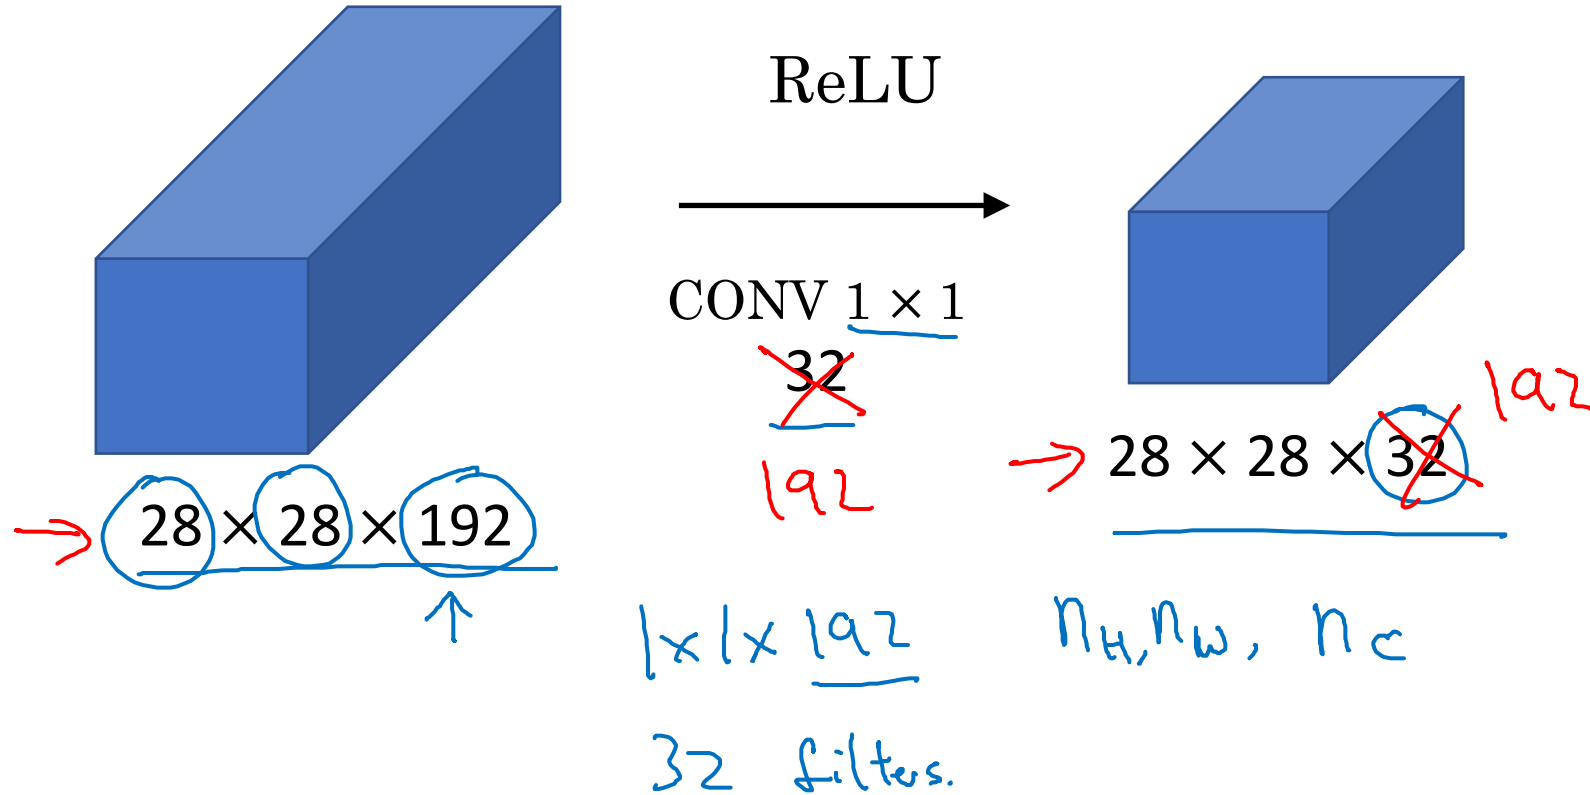
Network in  
Network

2	4	6	...		



$6 \times 6 \times \# \text{ filters}$

# Using $1 \times 1$ convolutions





Bahot tarah ke filters use kar lena ek hi time pe like ki see neech diagram. Let the network learn what filters it need. But the problem here is computation



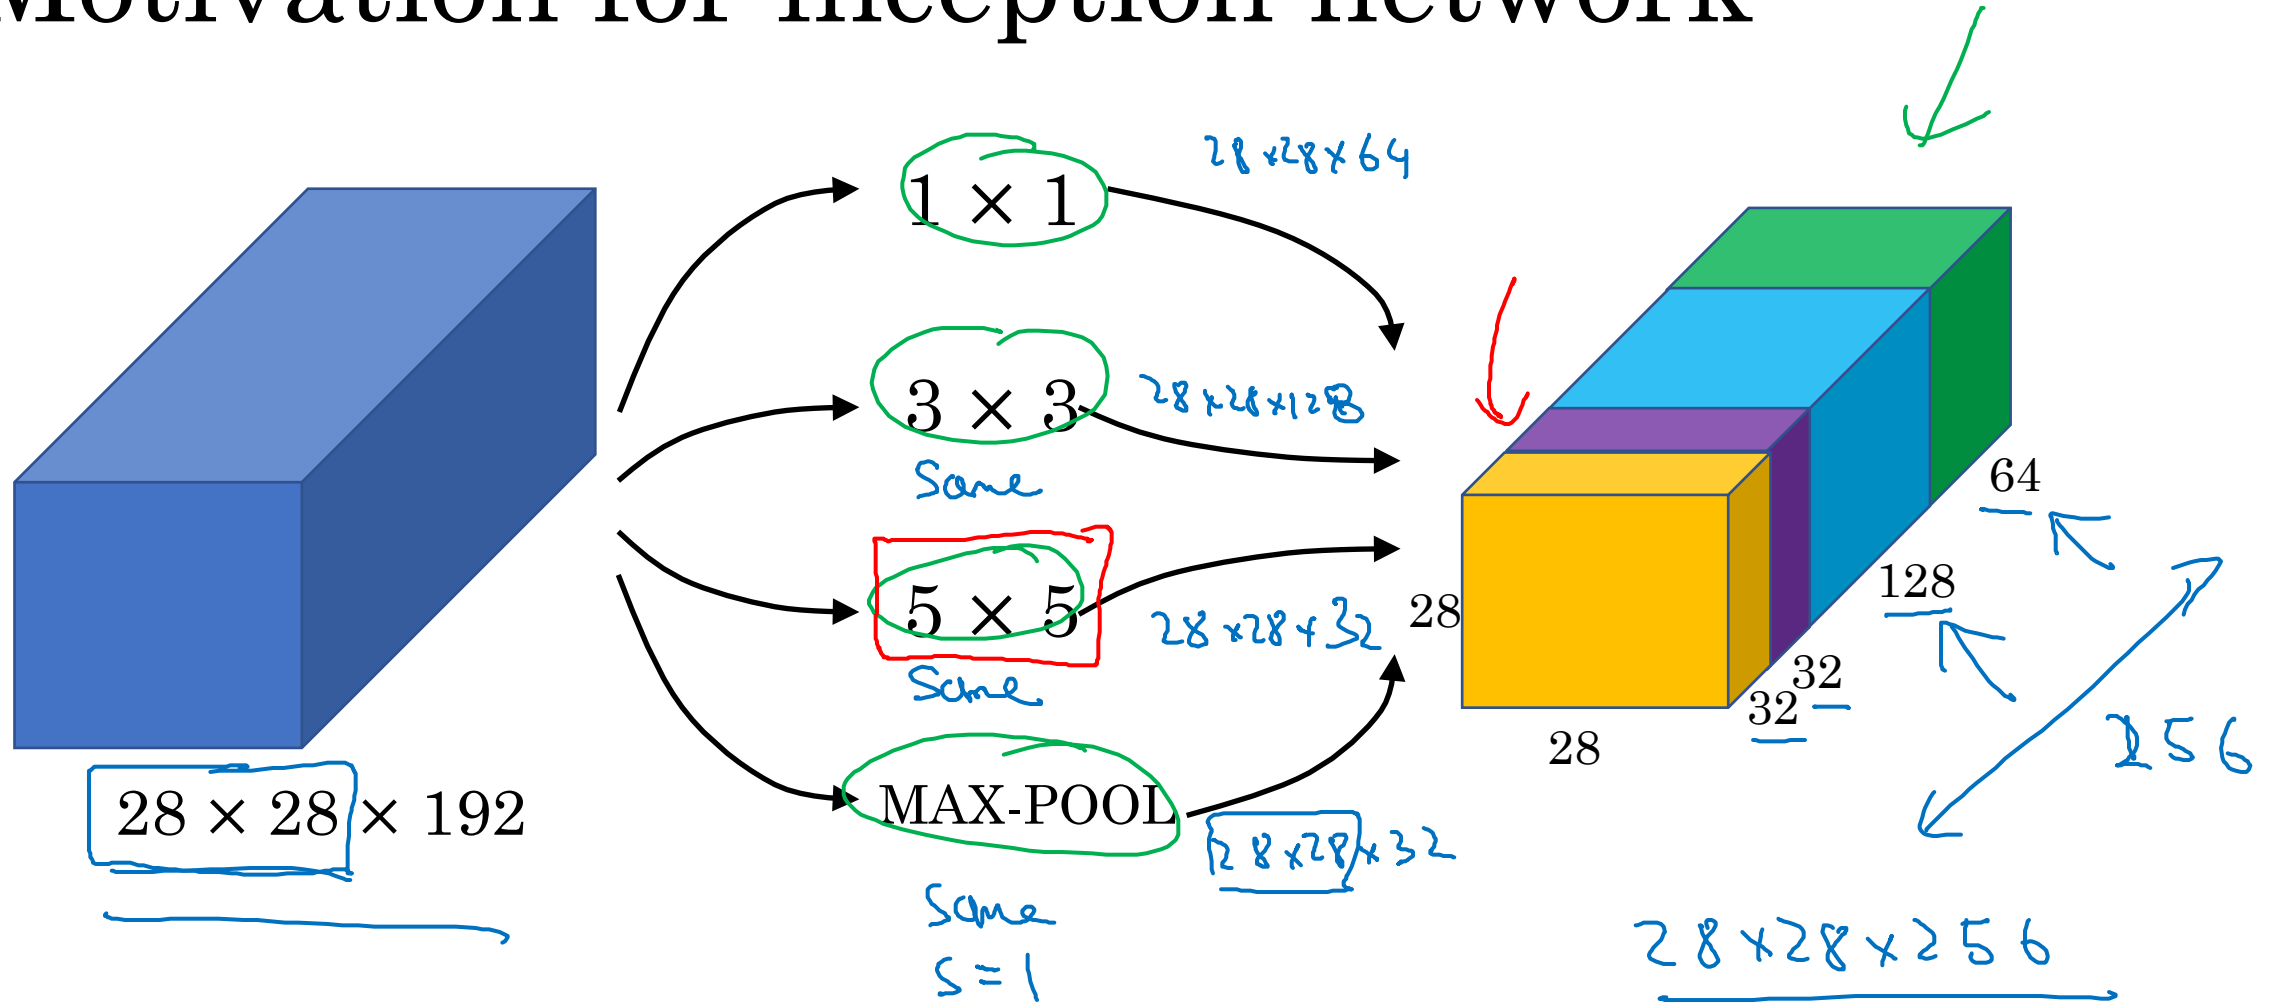
deeplearning.ai

# Case Studies


---

## Inception network motivation

# Motivation for inception network




# The problem of computational cost




# Using $1 \times 1$ convolution

"bottleneck layers"



The same above architecture can be achieved through  $1 \times 1$  filter and with much





deeplearning.ai

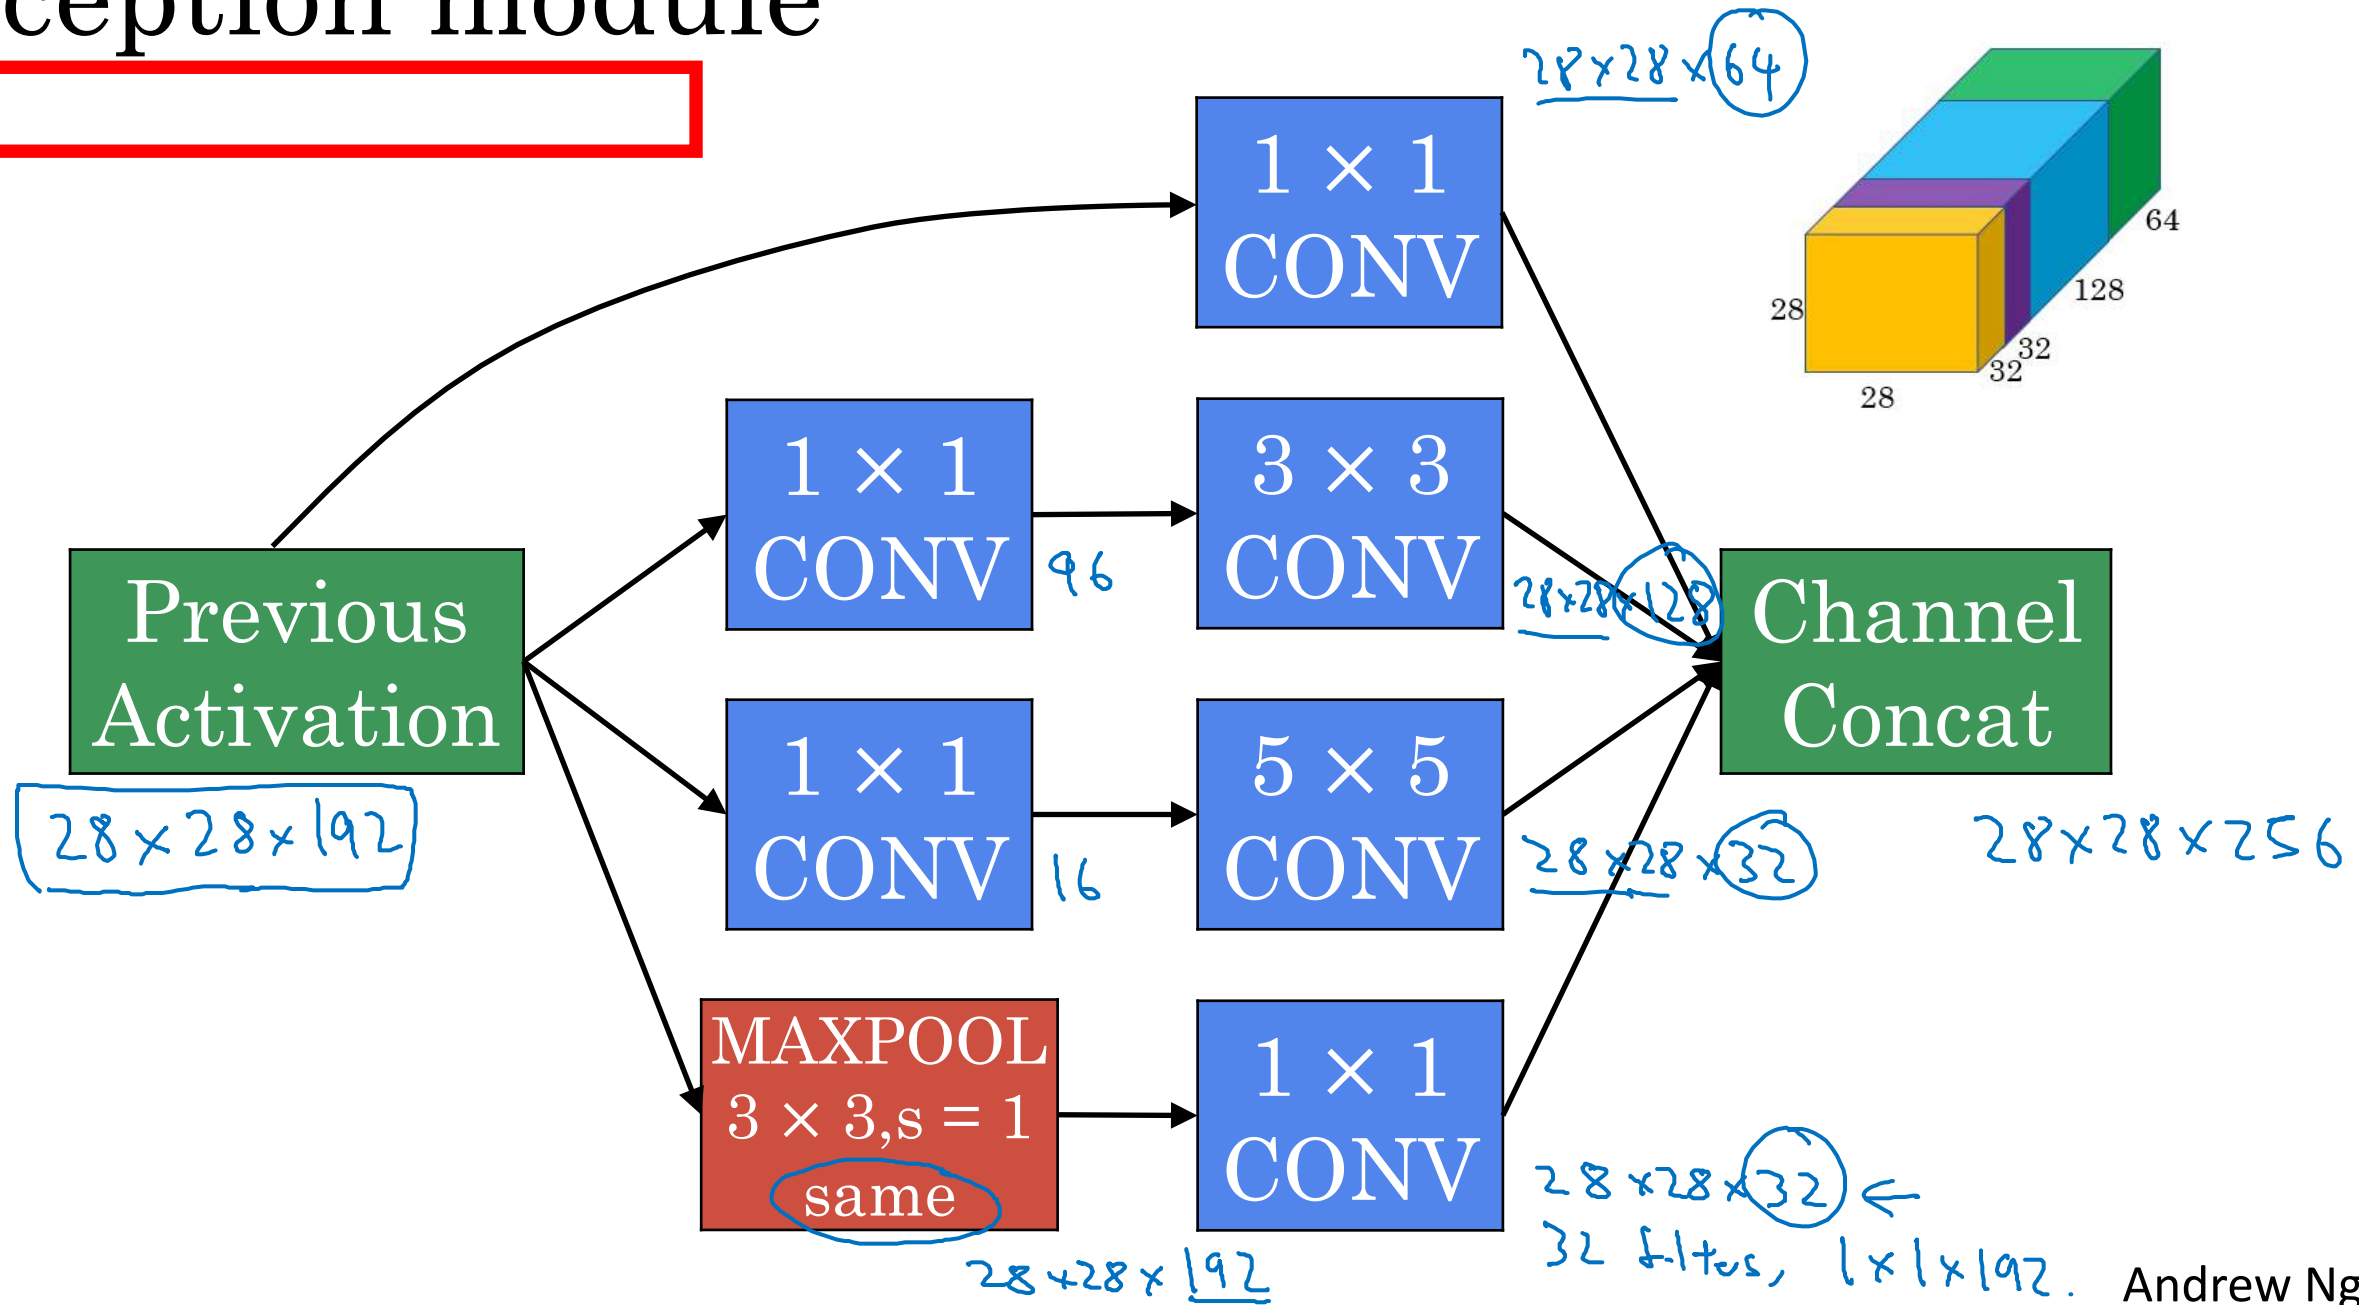
# Case Studies

---

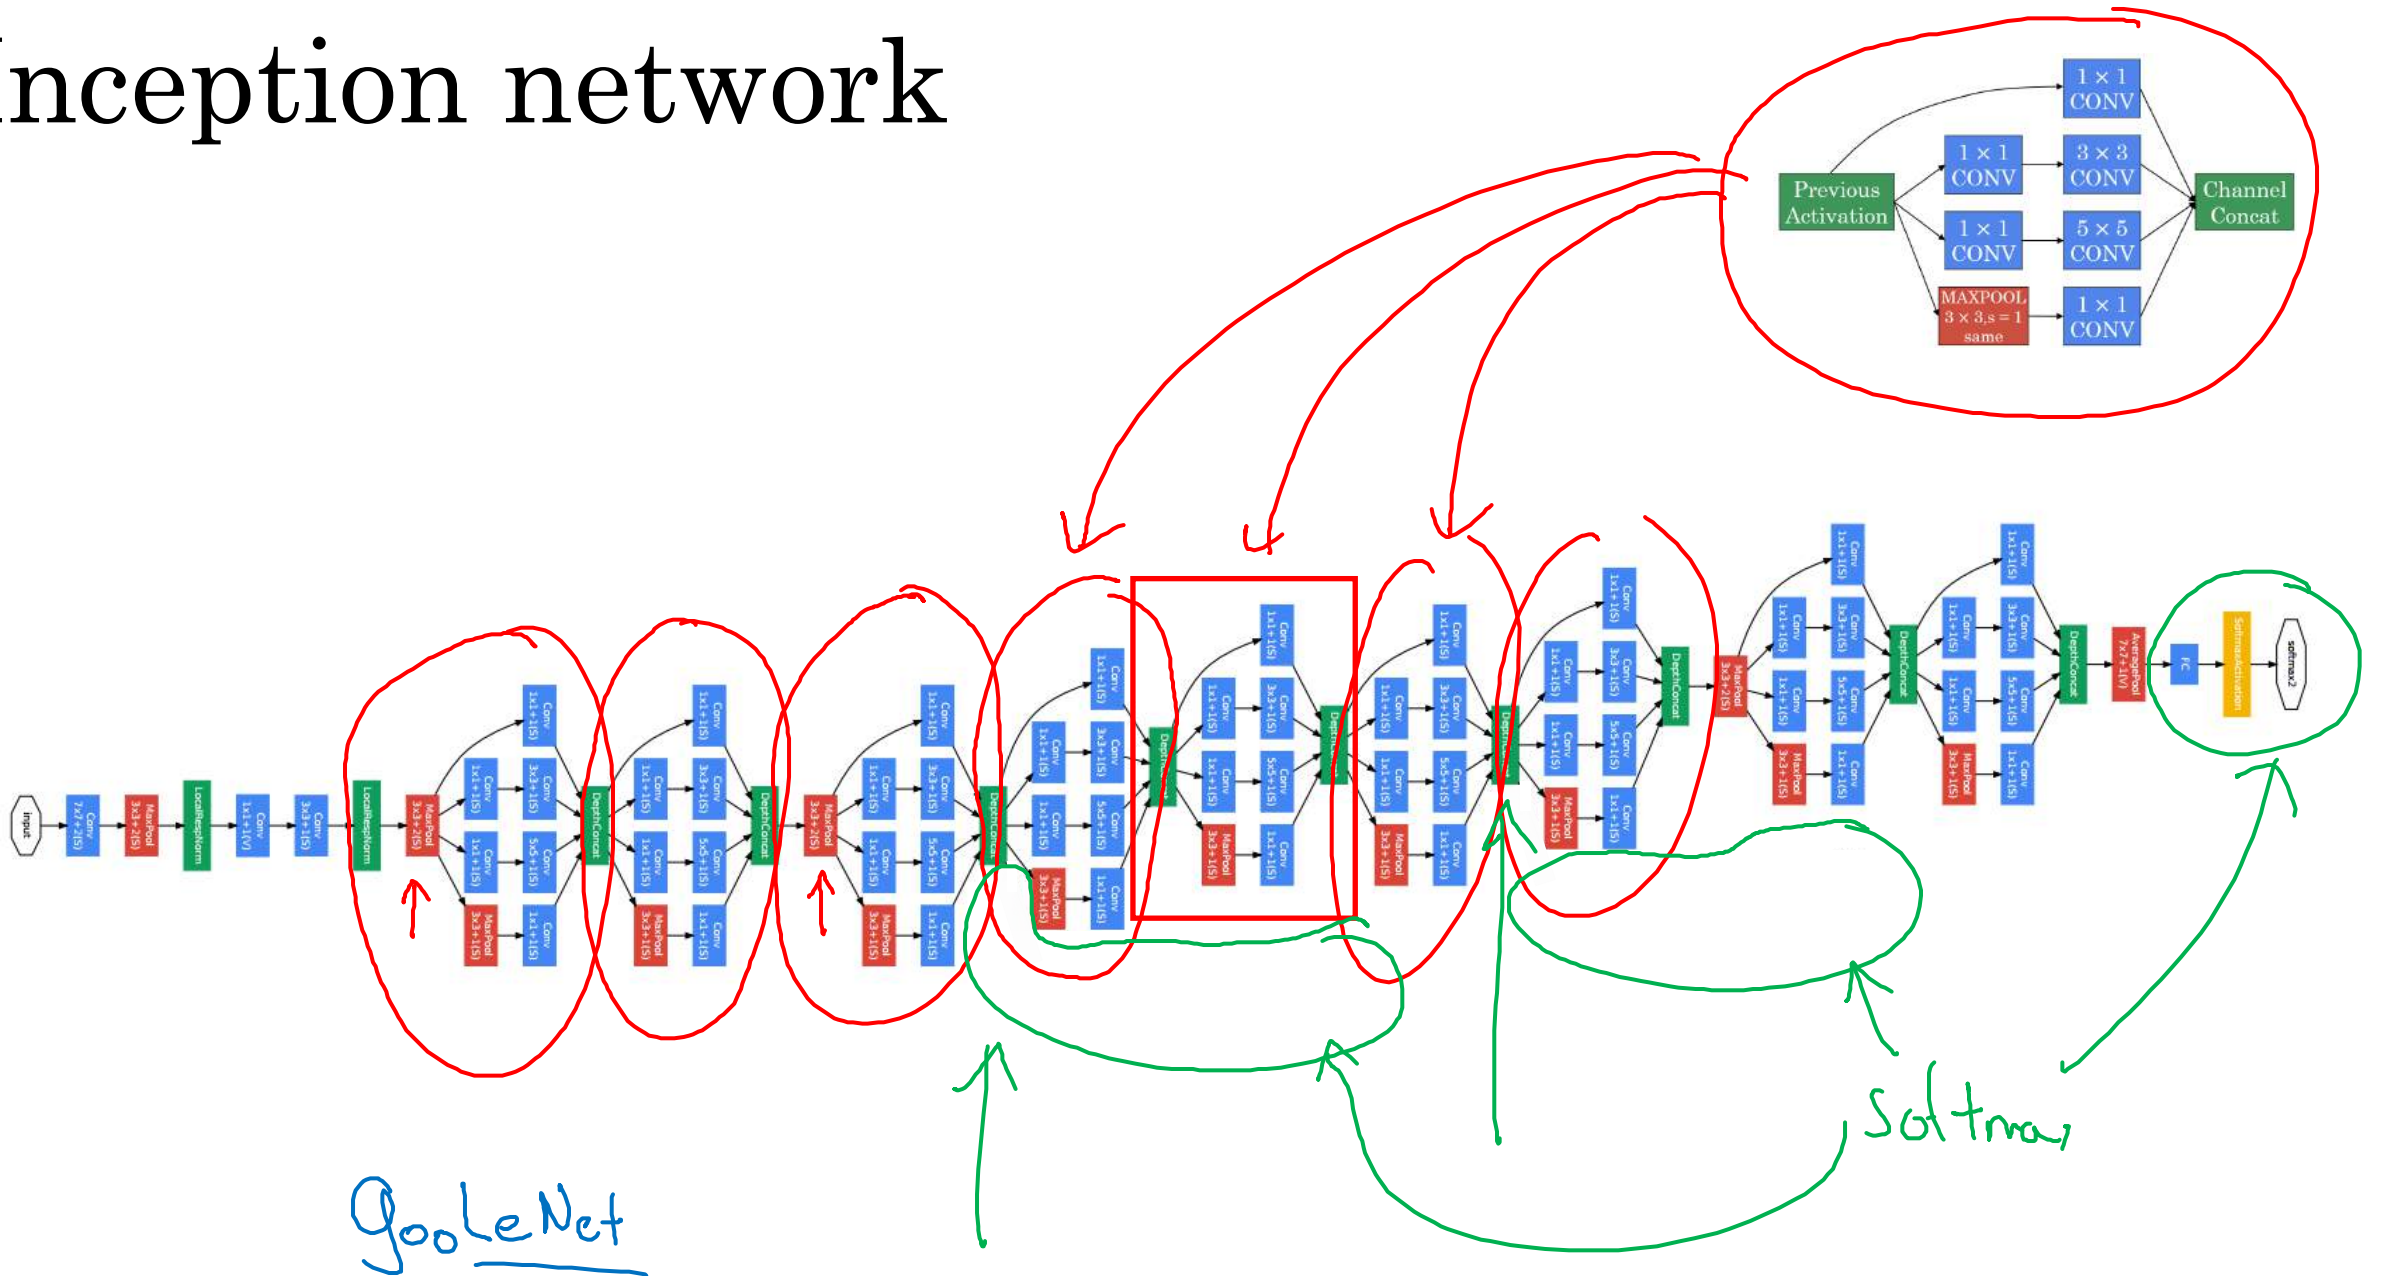
## Inception network

Below diagram representing how we can use inception in the diagram example taken above slide

# Inception module



# Inception network









deeplearning.ai

# Convolutional Neural Networks

---

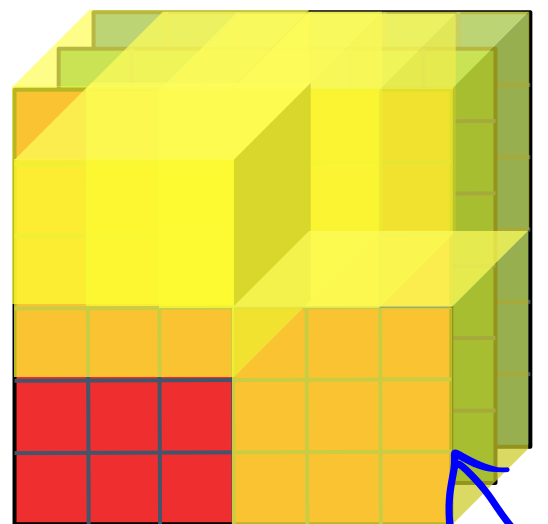
## MobileNet

# Motivation for MobileNets

- Low computational cost at deployment
- Useful for mobile and embedded vision applications
- Key idea: Normal vs. depthwise-separable convolutions

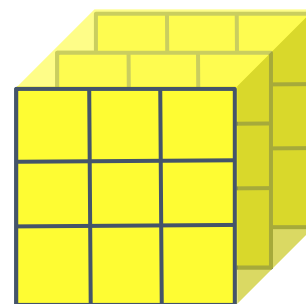


# Normal Convolution




$6 \times 6 \times 3$   
 $n \times n \times n_c$

\*



$3 \times 3 \times 3$   
 $f \times f \times n_c$

=

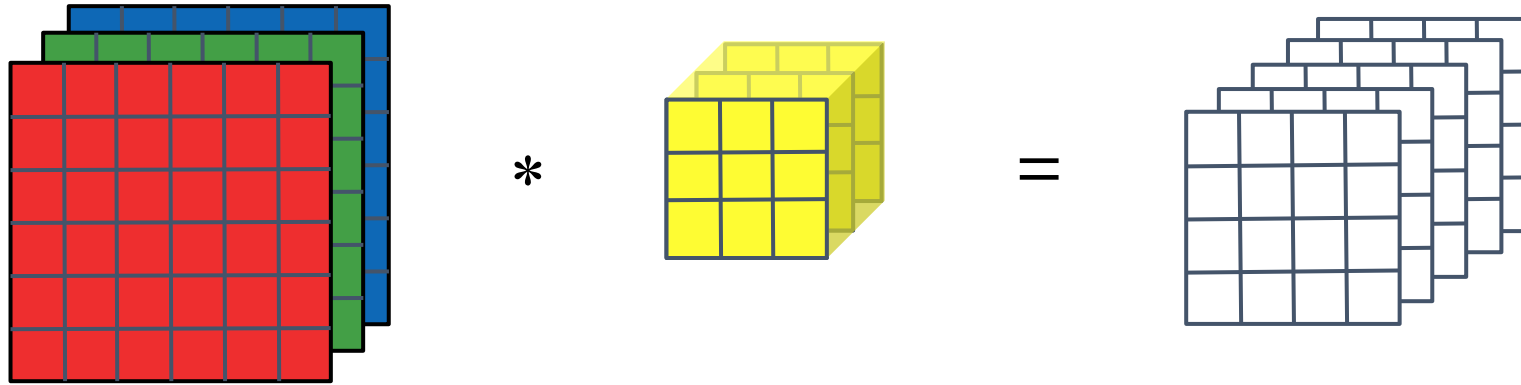


$4 \times 4 \times 5$   
 $n_{out} \times n_{out} \times n_c'$

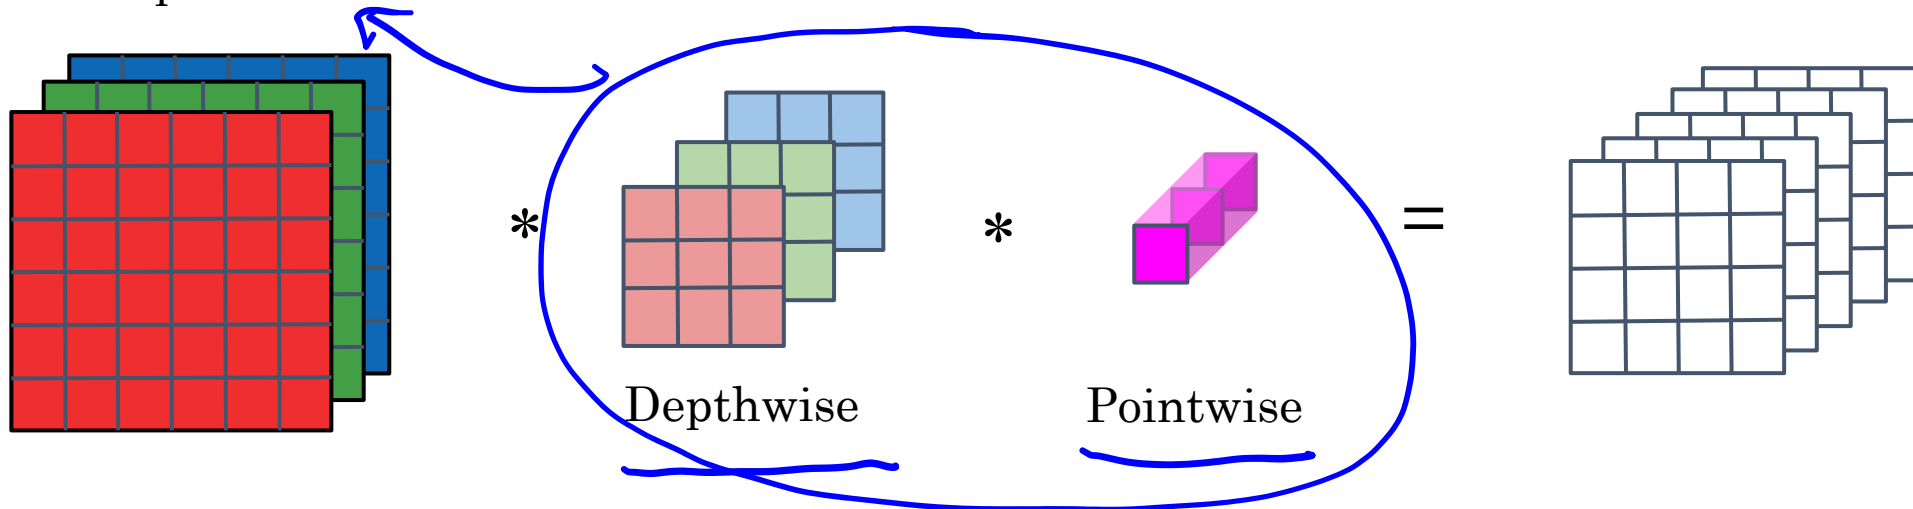
$$\begin{array}{ccccccc}
 \text{Computational cost} & = & \frac{\text{\#filter params}}{3 \times 3 \times 3} & \times & \frac{\text{\# filter positions}}{4 \times 4} & \times & \text{\# of filters} \\
 \rightarrow \underline{2160} & = & \uparrow & & \uparrow & & \uparrow \\
 & & & & & & 5
 \end{array}$$

# Depthwise Separable Convolution

Normal Convolution




Depthwise Separable Convolution



Here we do channel wise convultion like in 2D, red convoves with 1st layer, green with second adnd blue with thired and so an and they give

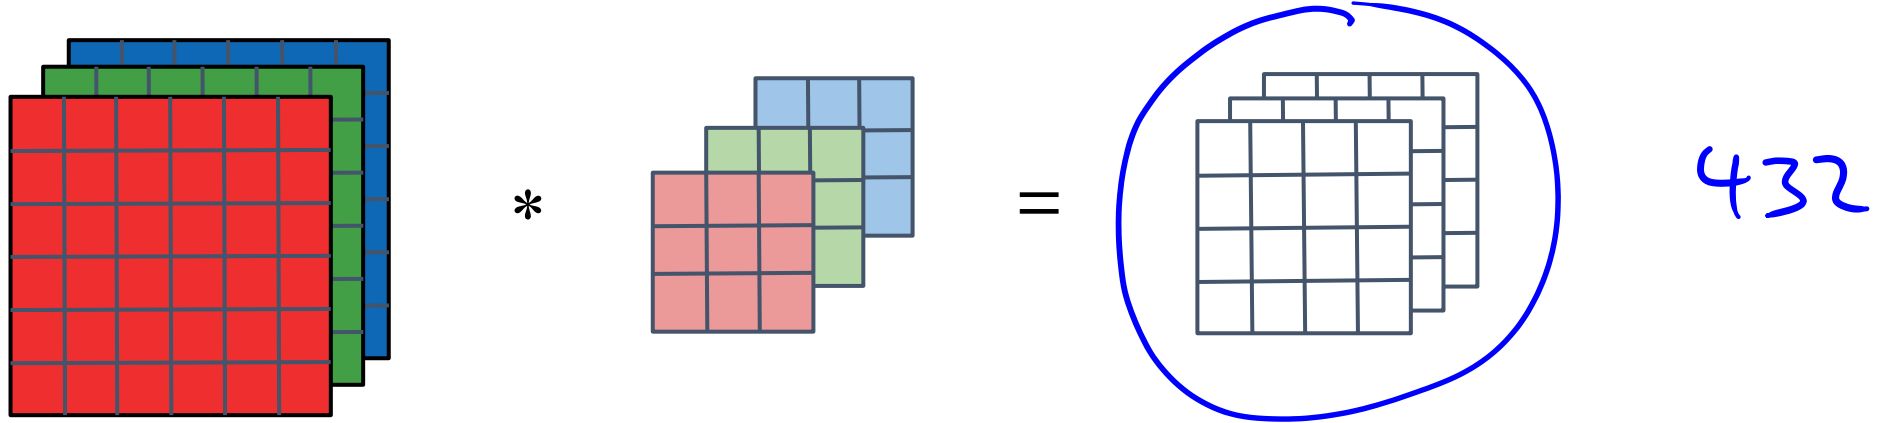
# Depthwise Convolution



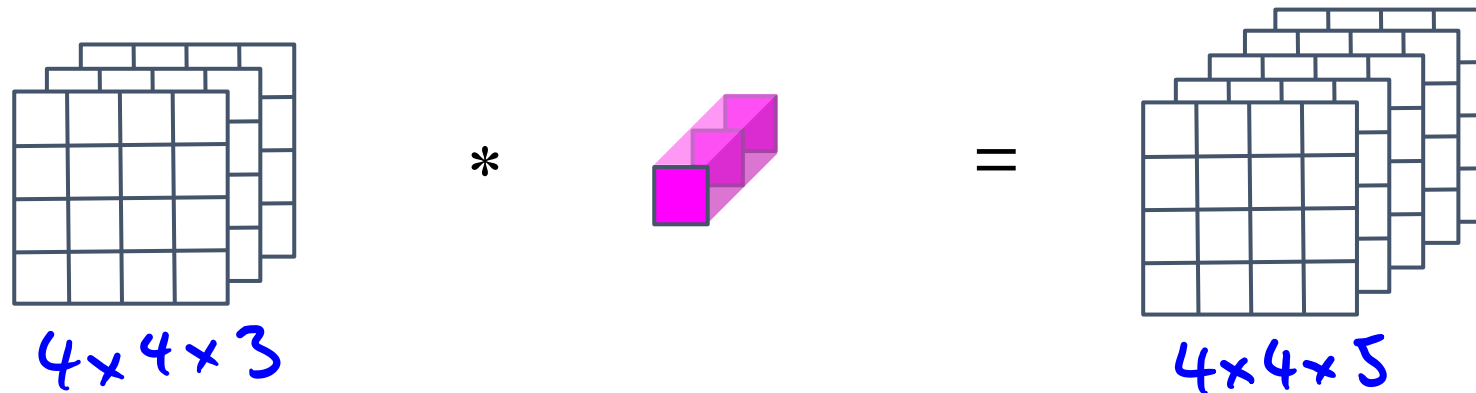
$$\begin{aligned}
 \text{Computational cost} &= \text{\#filter params} \times \text{\# filter positions} \times \text{\# of filters} \\
 432 &= \underbrace{3 \times 3} \times \underbrace{4 \times 4 \times 3}
 \end{aligned}$$

# Depthwise Separable Convolution

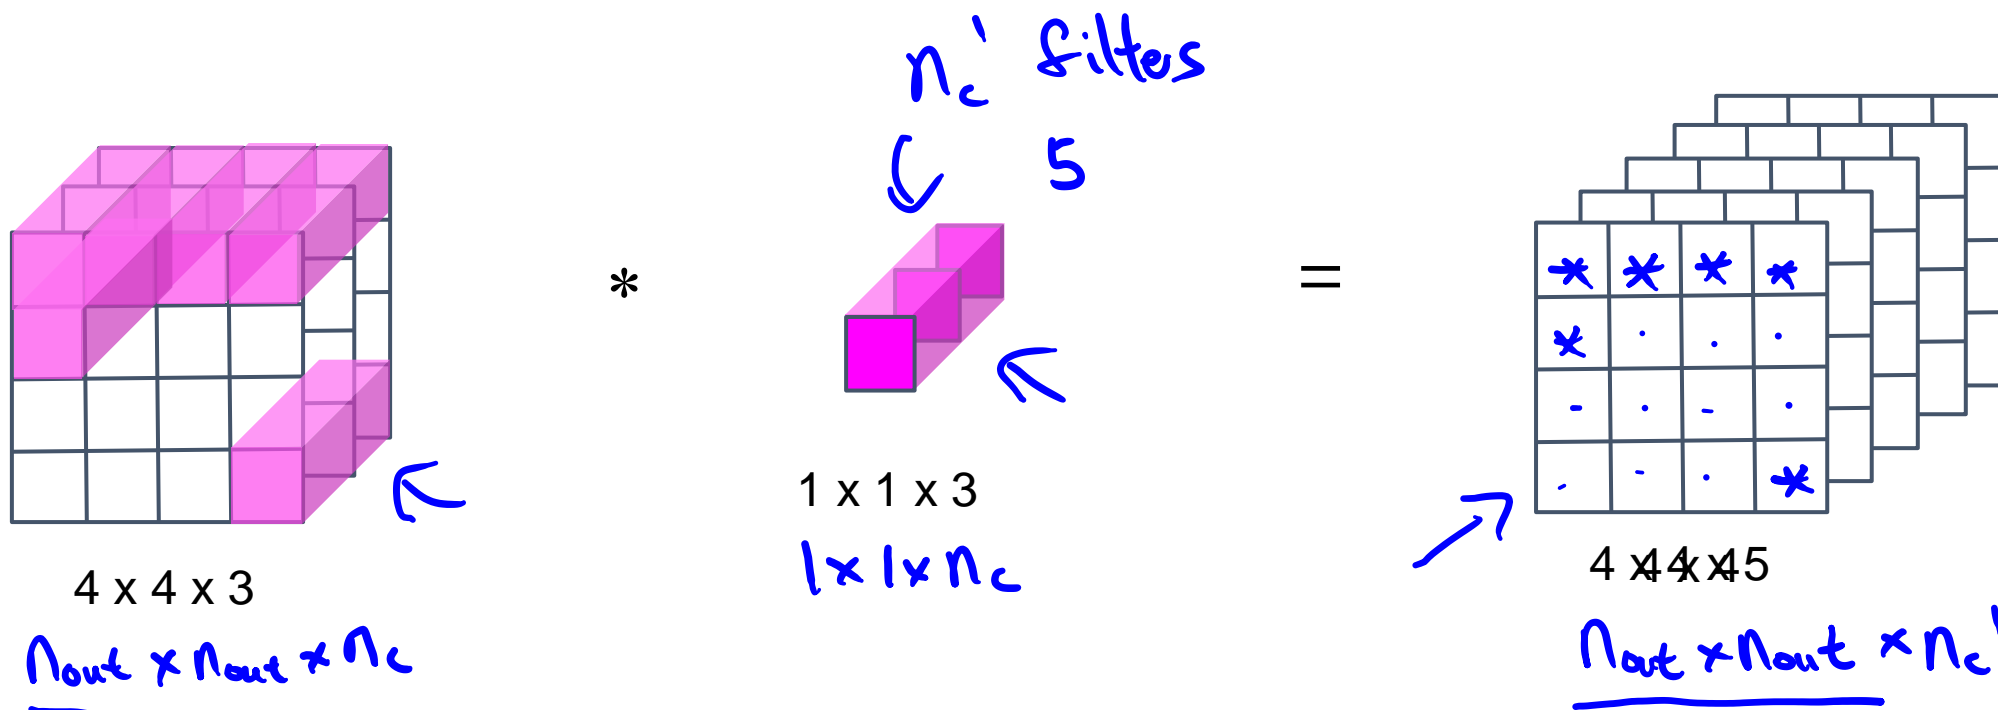
Depthwise Convolution



Pointwise Convolution



# Pointwise Convolution

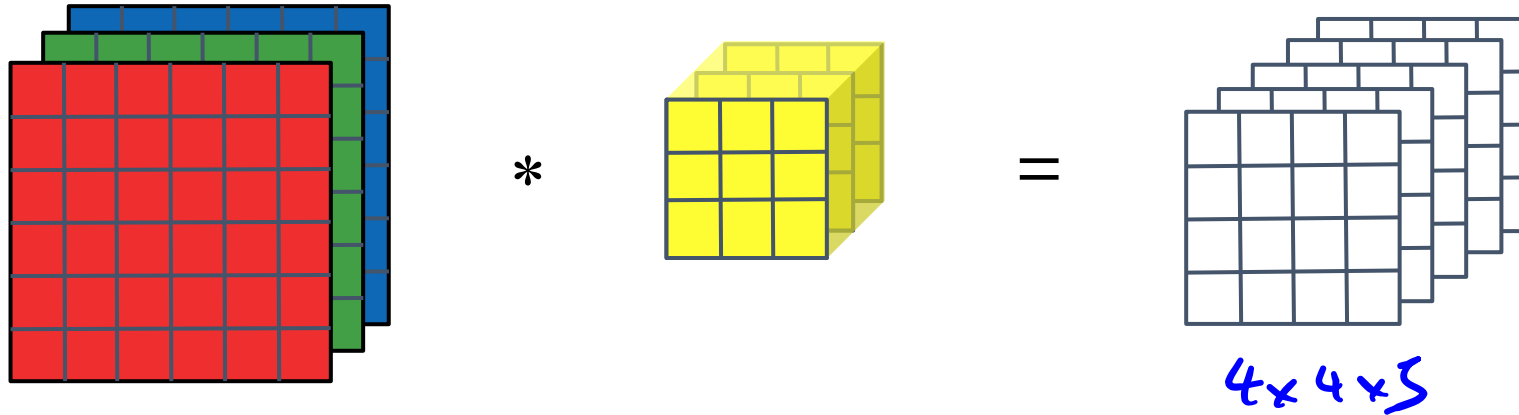


Computational cost = #filter params x # filter positions x # of filters

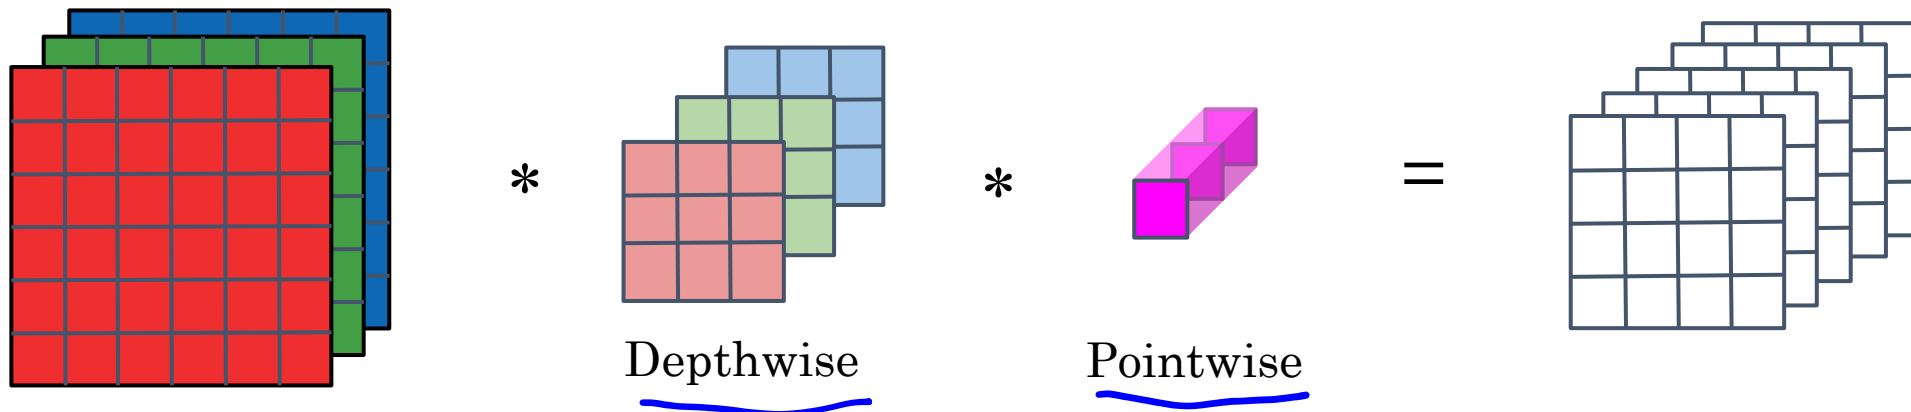
240 =  $1 \times 1 \times 3$  x  $4 \times 4$  x 5

# Depthwise Separable Convolution

Normal Convolution



Depthwise Separable Convolution





# Cost Summary

Cost of normal convolution

2160

Cost of depthwise separable convolution

$$\begin{array}{ccc} \text{depthwise} & + & \text{pointwise} \\ 432 & + & 240 = 672 \end{array}$$

$$\frac{672}{2160} = 0.31 \leftarrow$$

$$= \frac{1}{\textcircled{n_c}} + \frac{1}{f^2}$$
$$\frac{1}{5} + \frac{1}{9}$$




$$= \frac{1}{512} + \frac{1}{3^2}$$

↑                      ↑




~10 times cheaper

# Depthwise Separable Convolution

Depthwise Convolution



Pointwise Convolution





deeplearning.ai

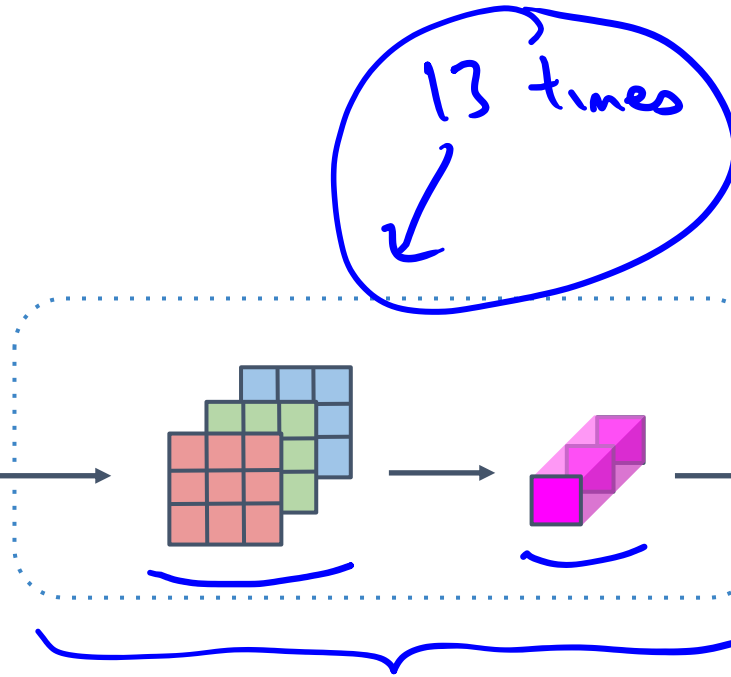
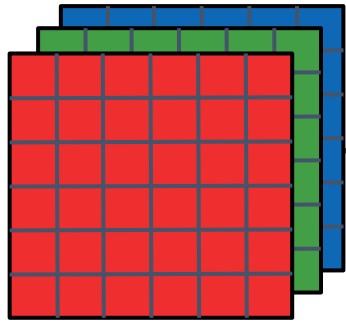
# Convolutional Neural Networks

---

## MobileNet Architecture

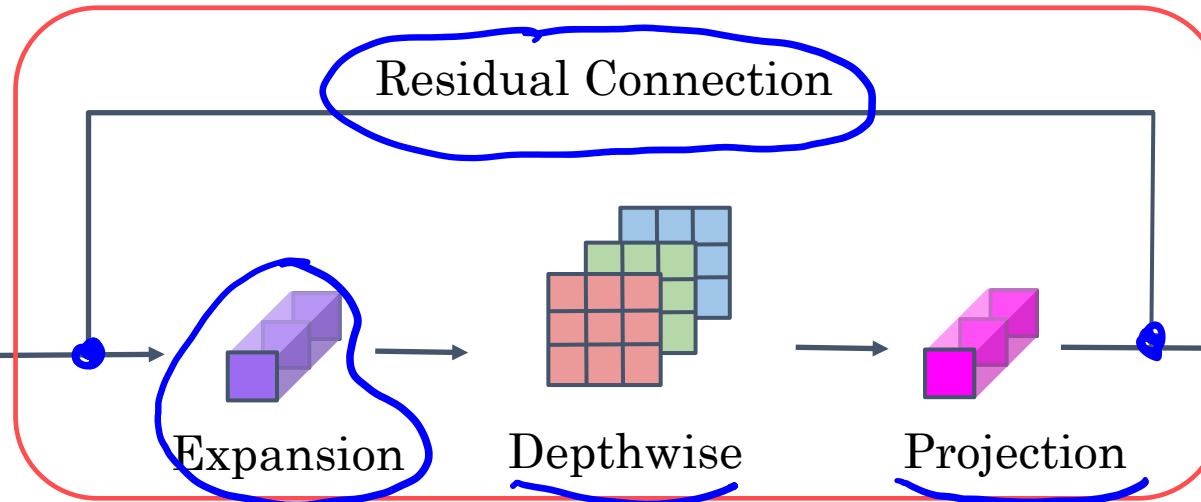
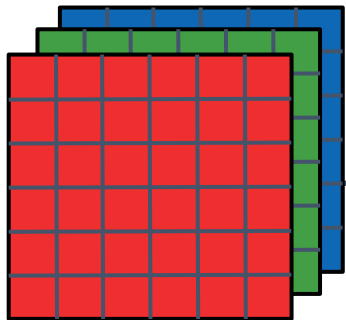
# MobileNet

MobileNet v1



Pool, FC, Softmax

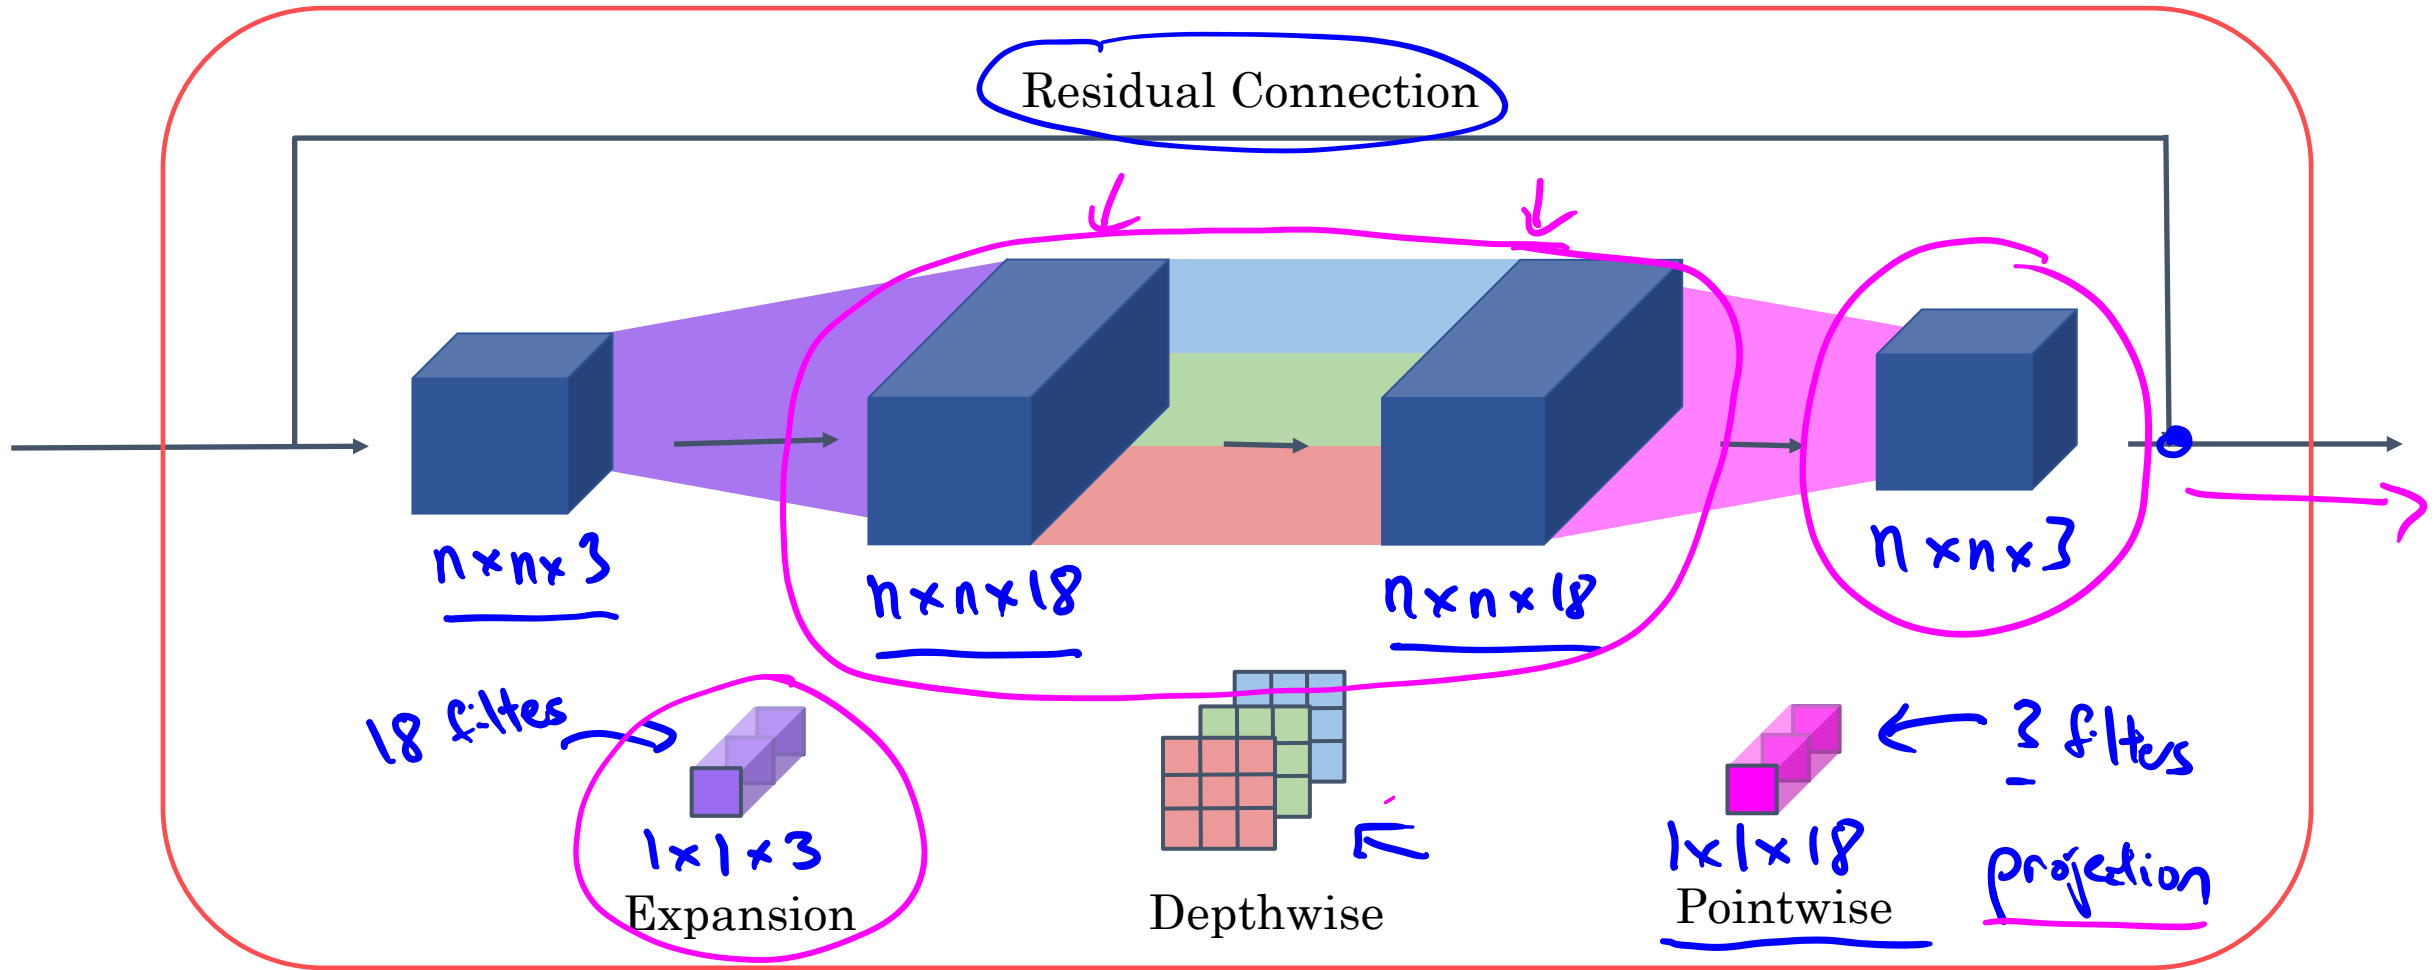
MobileNet v2



Pool, FC, Softmax


Bottleneck

# MobileNet v2 Bottleneck

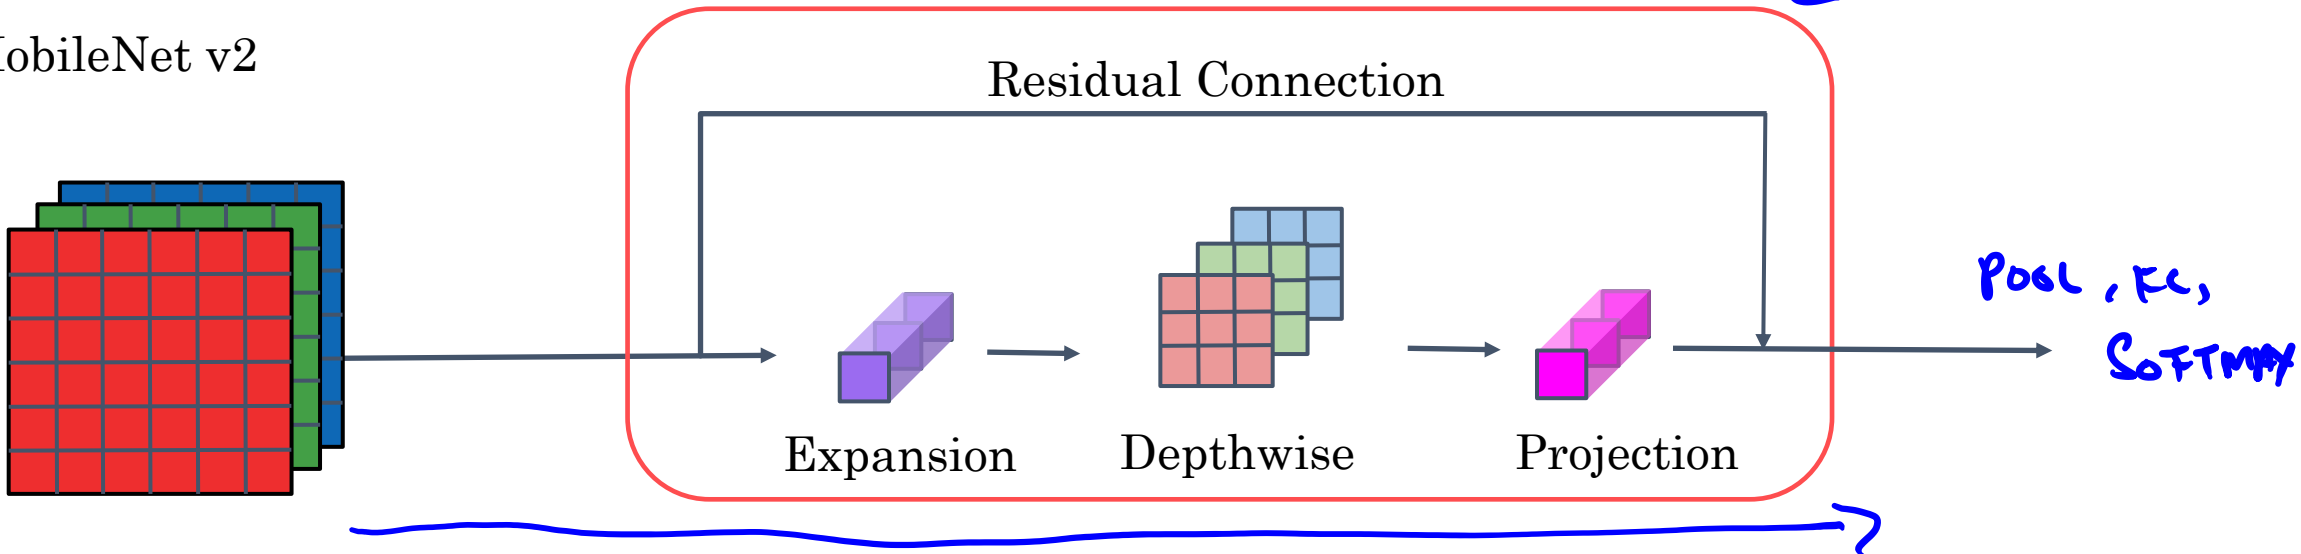


# MobileNet

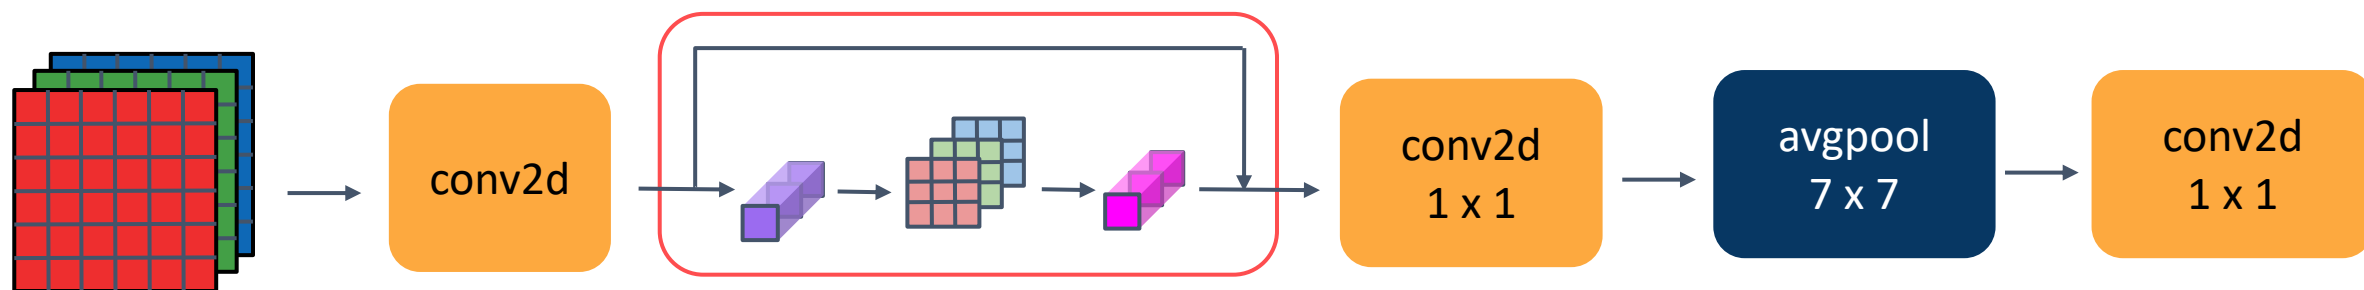
MobileNet v1



MobileNet v2



# MobileNet v2 Full Architecture





deeplearning.ai

# Convolutional Neural Networks


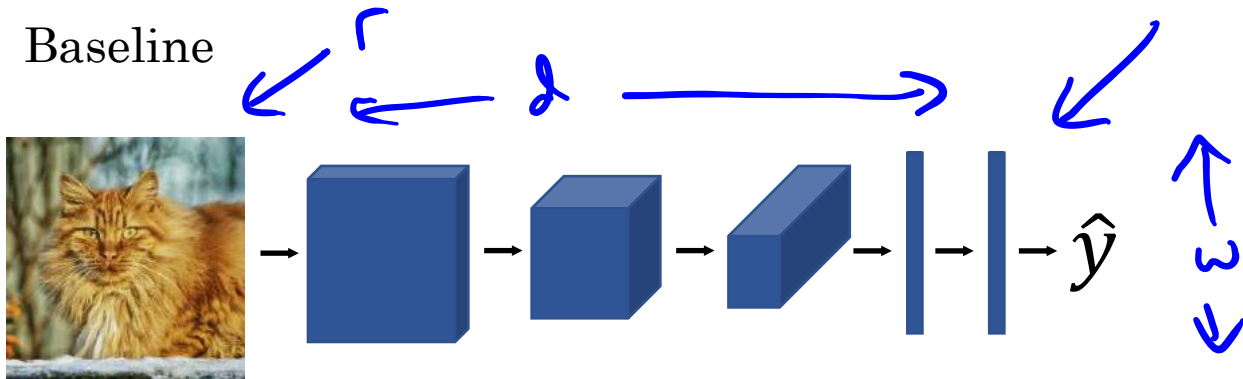
---

## EfficientNet



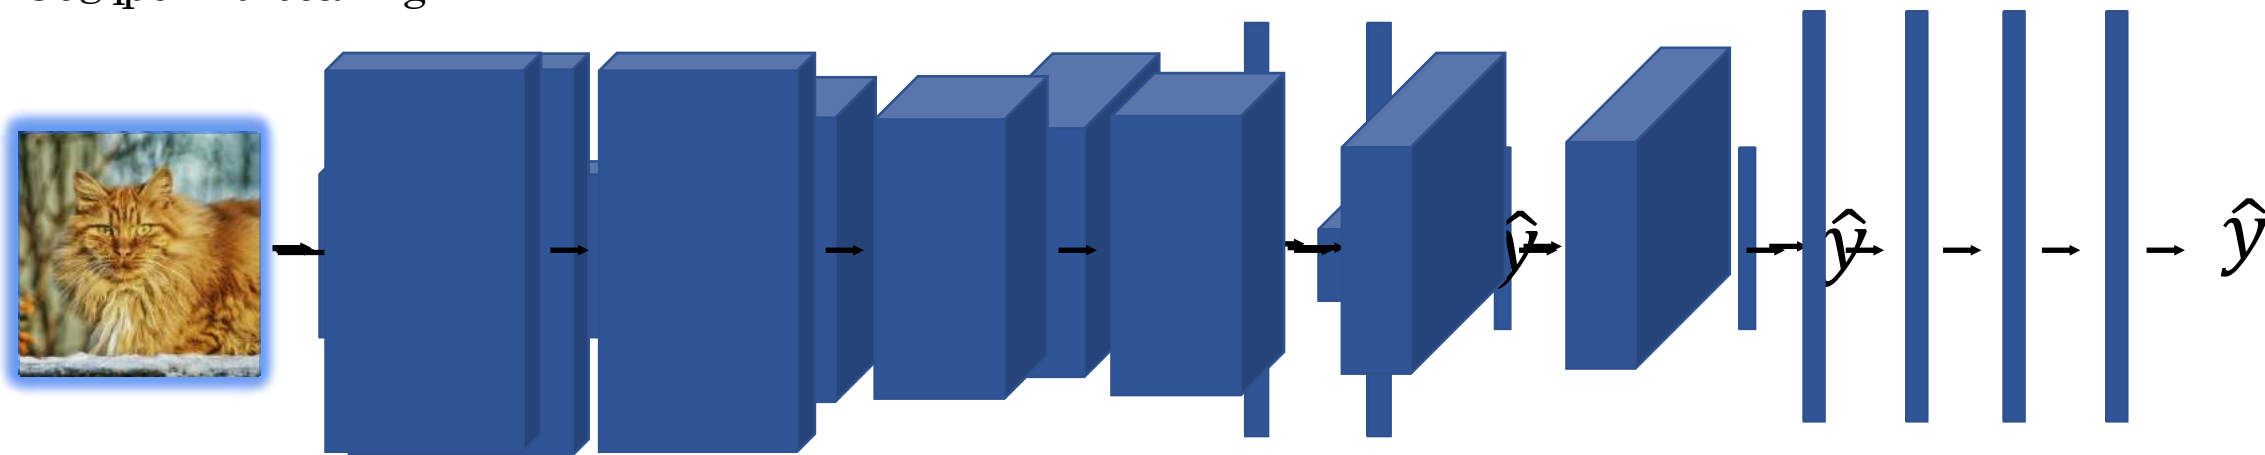
# EfficientNet

## Baseline



A diagram showing a vertical oval representing a 3D volume. Three arrows point from the left towards the oval, each aligned with one of its internal labels: 'r' at the top, 'd' in the middle, and 'w' at the bottom. To the right of the oval, the words 'resolution', 'depth', and 'width' are written, corresponding to 'r', 'd', and 'w' respectively.

# Wider Rescating





deeplearning.ai

# Practical advice for using ConvNets


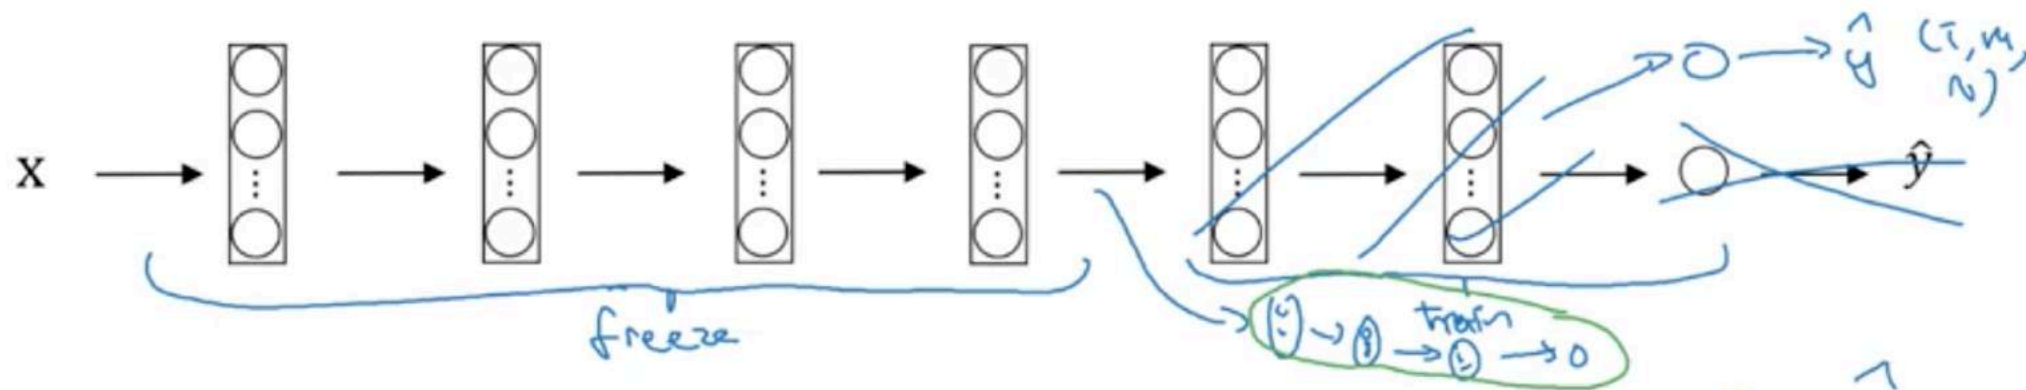

---

## Transfer Learning

Instead of randomly initialising wts and building things from scratch , we can use wts that have already

# Transfer Learning

 Tigger
 Misty
Neither





deeplearning.ai

# Practical advice for using ConvNets

---

## Data augmentation


# Common augmentation method

## Mirroring




y

## Random Cropping




- Rotation
- Shearing
- Local warping
- ...





# Color shifting



Advanced:

PCA

[ml-class.org](http://ml-class.org)

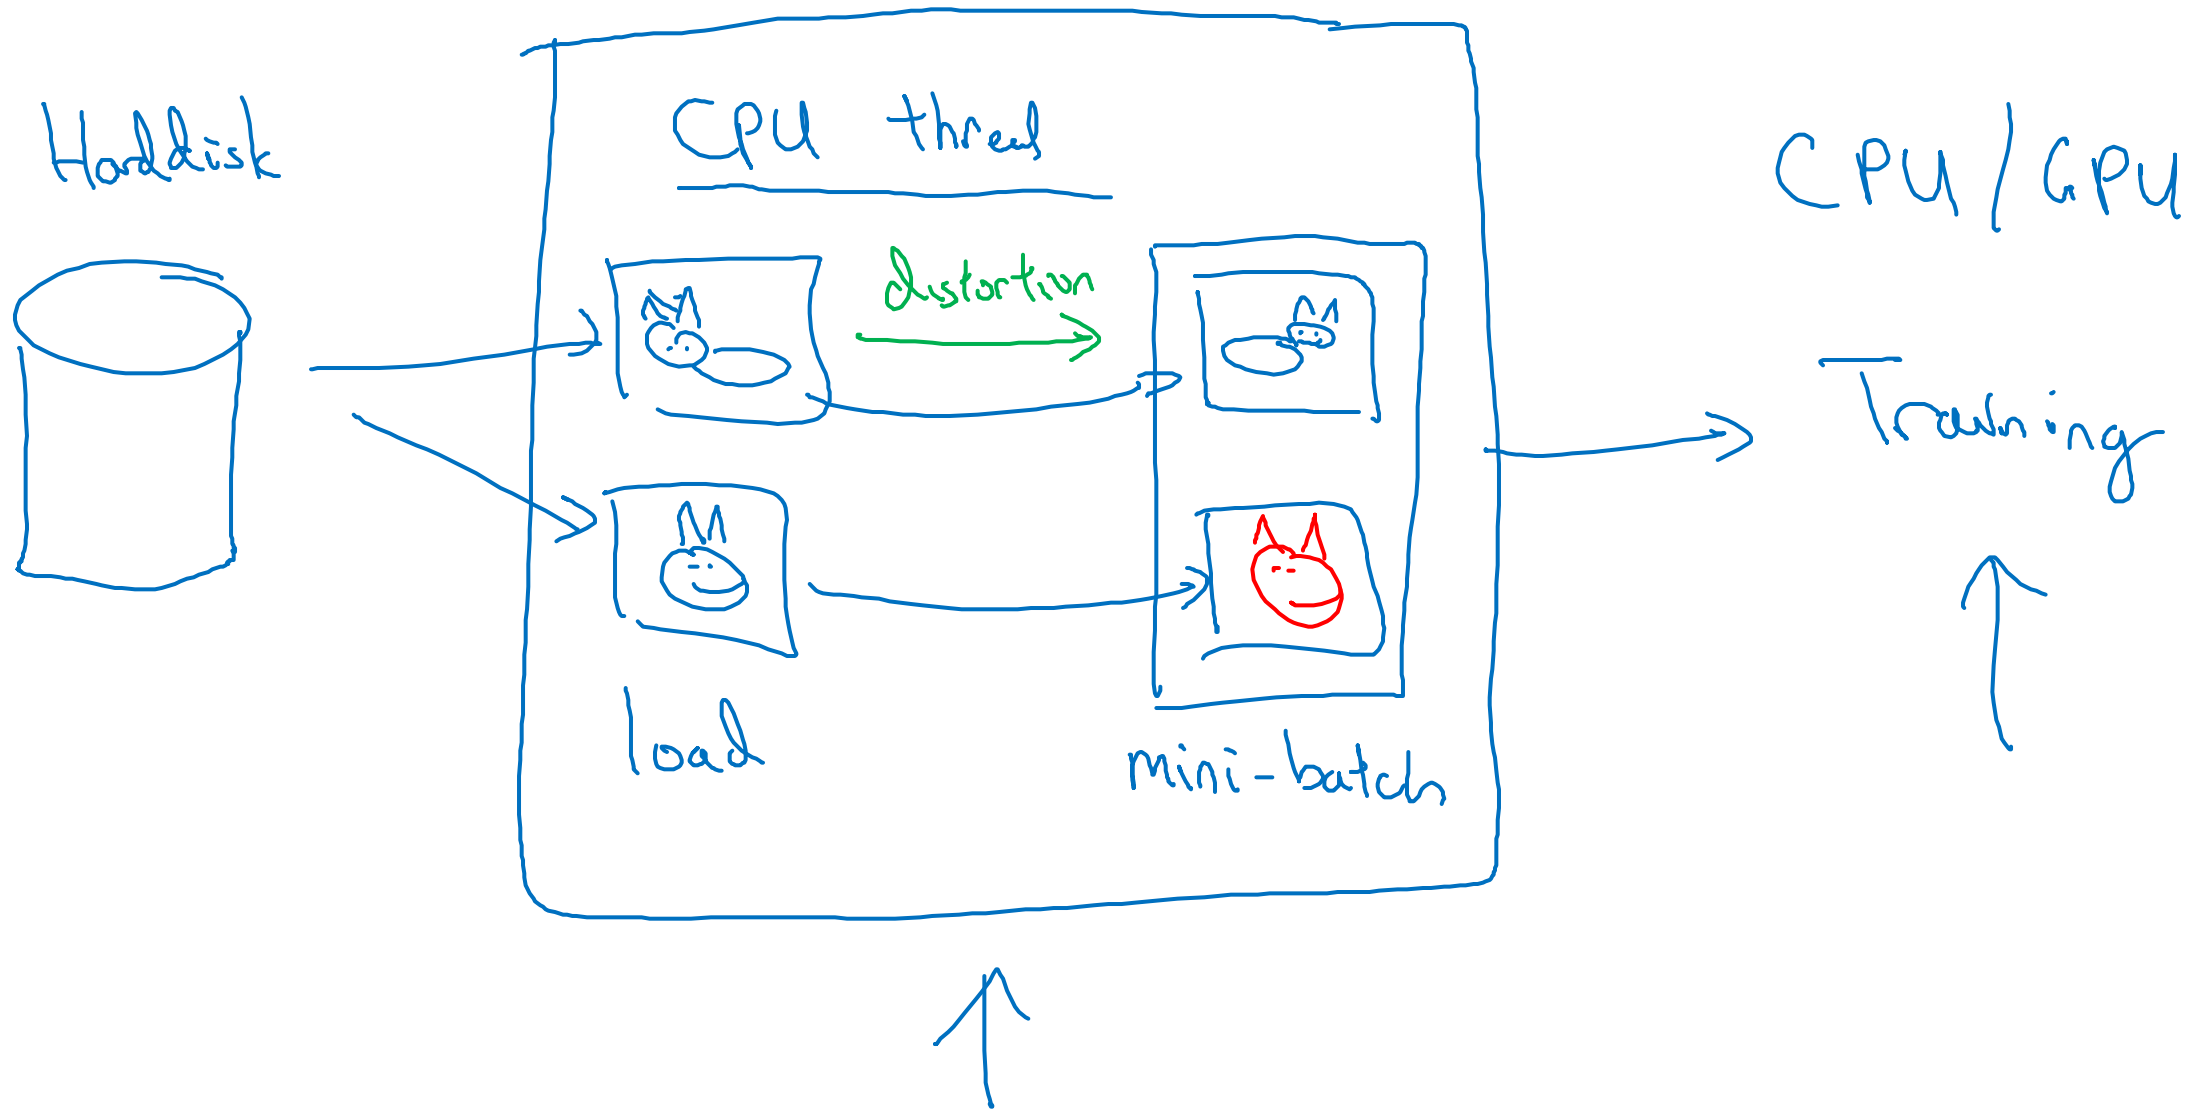
[ AlexNet paper

[ "PCA color augmentation."

R B

G

# Implementing distortions during training





deeplearning.ai

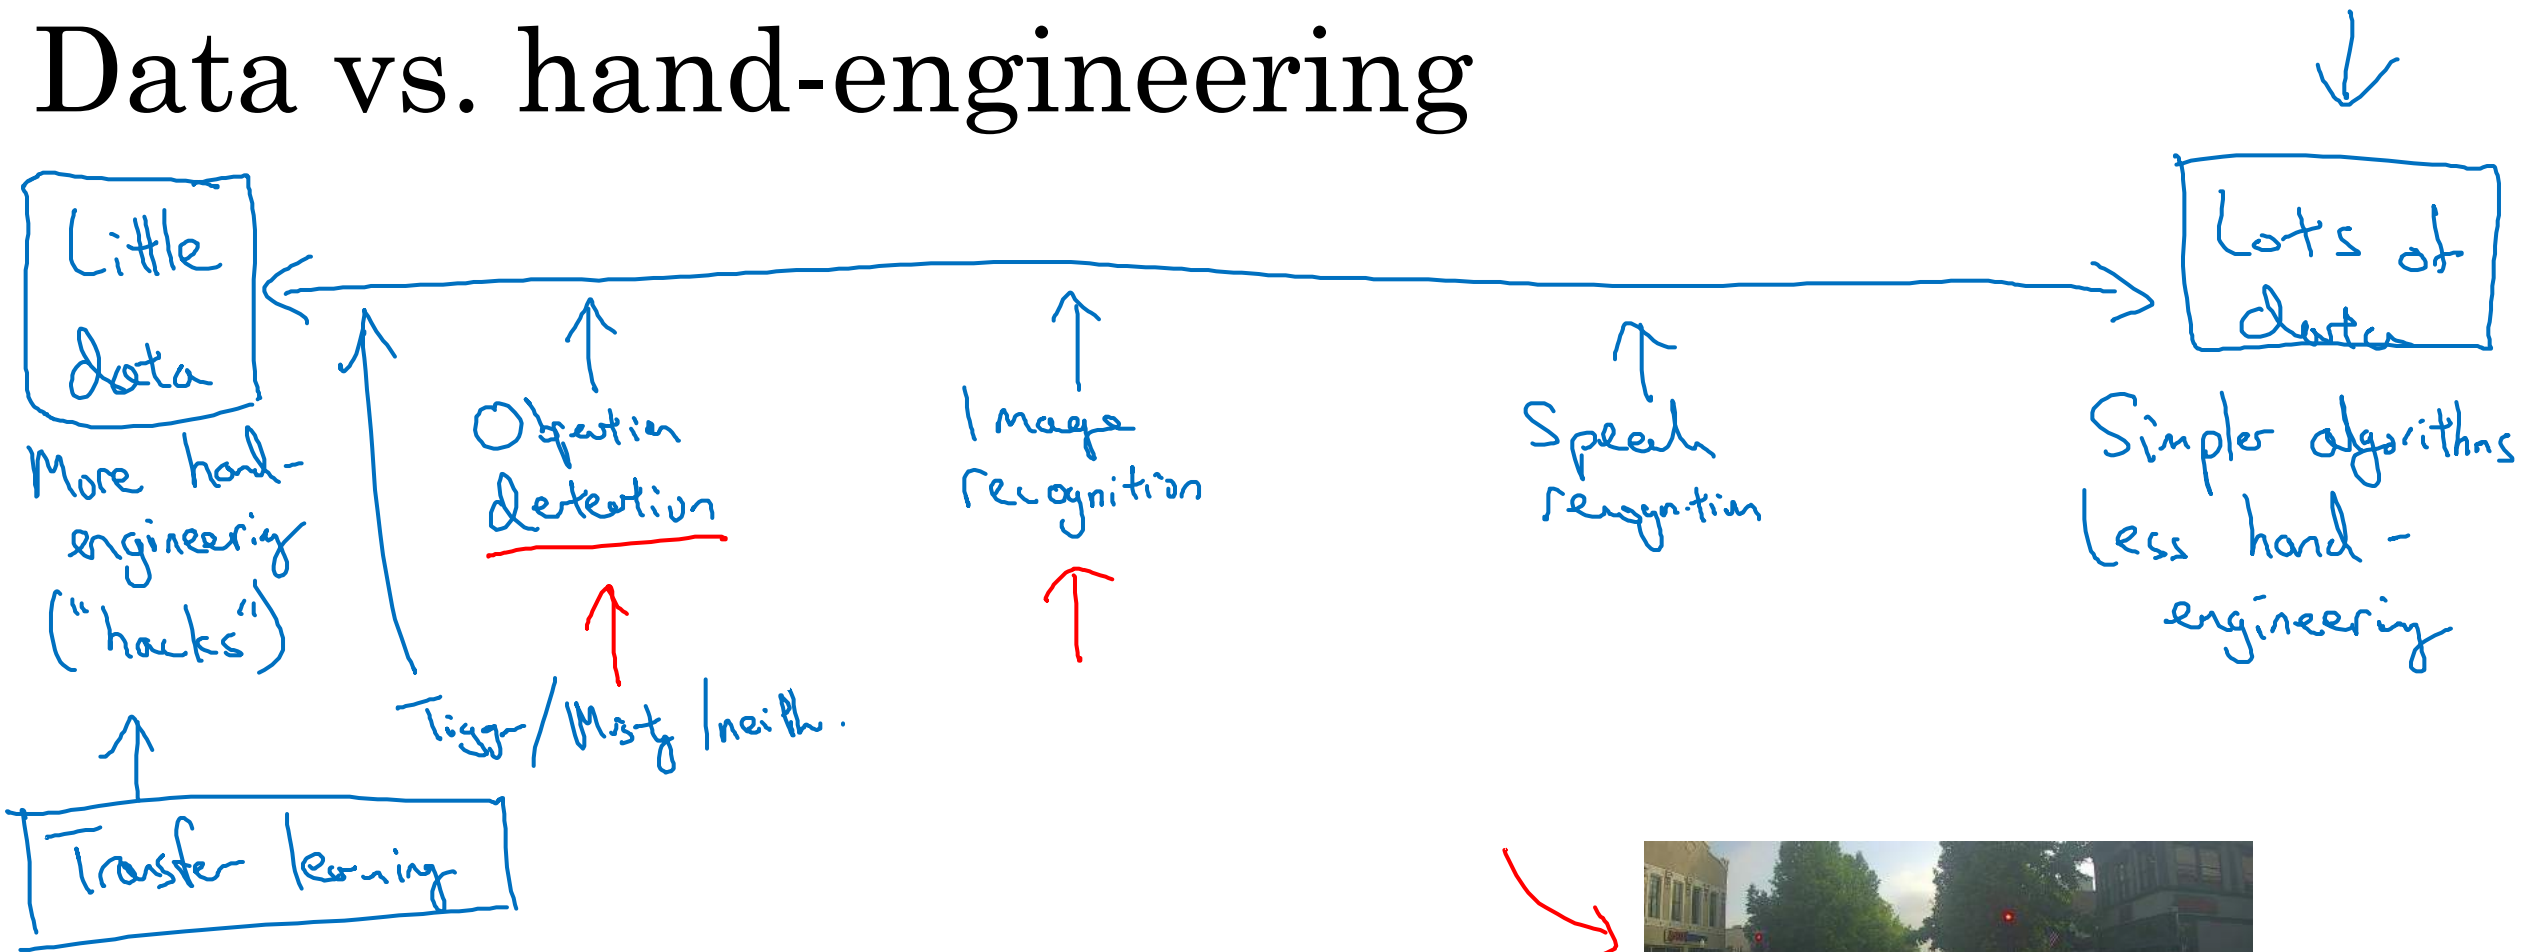
# Practical advice for using ConvNets

---

## The state of computer vision



# Data vs. hand-engineering



Two sources of knowledge

- • Labeled data  $(x, y)$
- • Hand engineered features/network architecture/other components



# Tips for doing well on benchmarks/winning competitions

## Ensembling

3-15 networks

→  $\hat{y}$

- Train several networks independently and average their outputs

## Multi-crop at test time

- Run classifier on multiple versions of test images and average results

10-crop



1

+



4

+



1

+



4

# Use open source code

- Use architectures of networks published in the literature
- Use open source implementations if possible
- Use pretrained models and fine-tune on your dataset